


# Virtual Memory



CS 351: Systems Programming  
Michael Lee <lee@iit.edu>




IIT College of Science  
ILLINOIS INSTITUTE OF TECHNOLOGY



previously: SRAM  $\Leftrightarrow$  DRAM





next: DRAM  $\Leftrightarrow$  HDD, SSD, etc.  
i.e., memory as a “cache” for disk



main goals:

1. maximize memory *throughput*
2. maximize memory *utilization*
3. provide *address space consistency*  
& *memory protection* to processes



*throughput* = # bytes per second

- depends on access latencies (DRAM, HDD, etc.) and “hit rate”



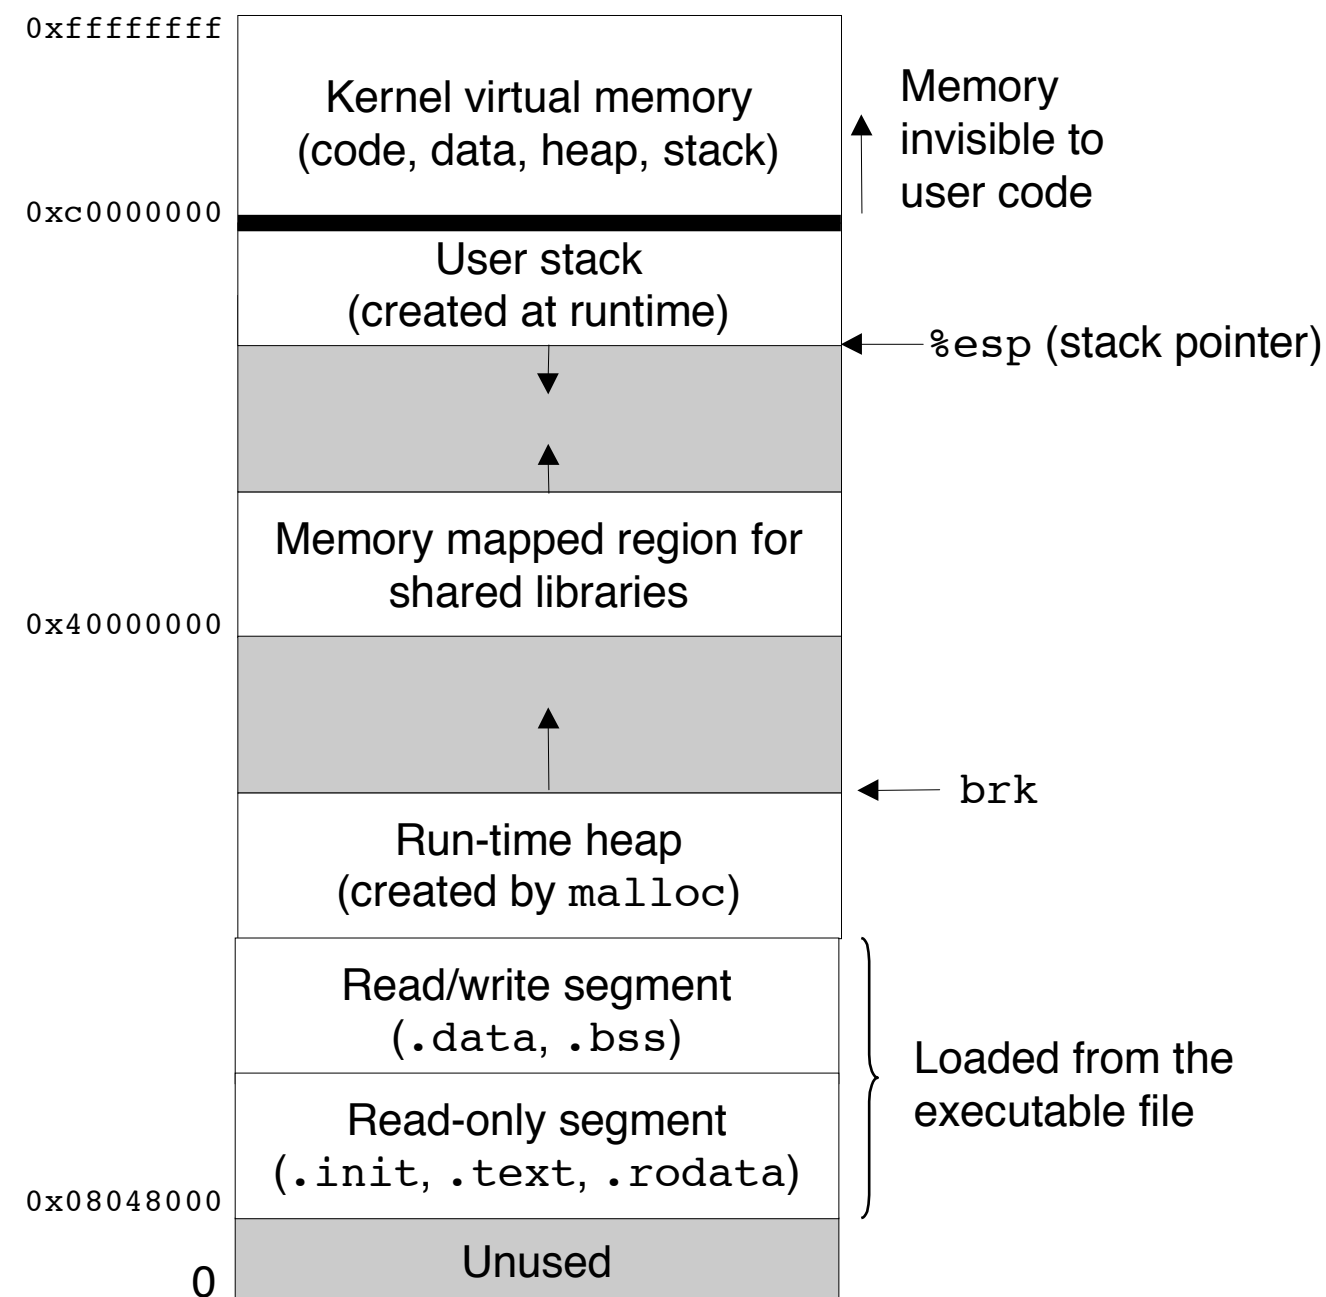
*utilization* = fraction of allocated memory that contains “user” data (aka *payload*)

- vs. metadata and other overhead required for memory allocation



*address space consistency* → provide a uniform  
“view” of memory to each process



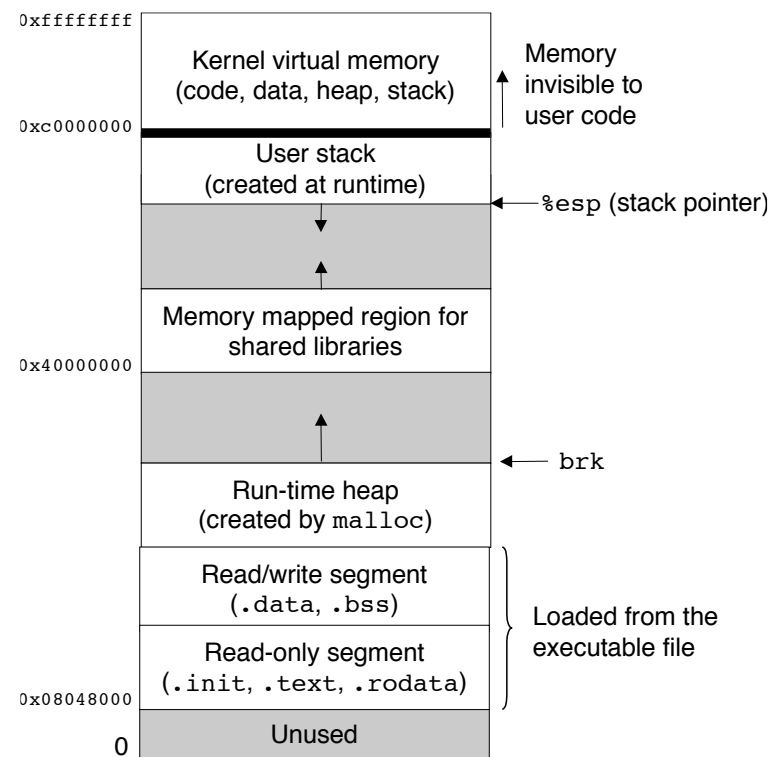
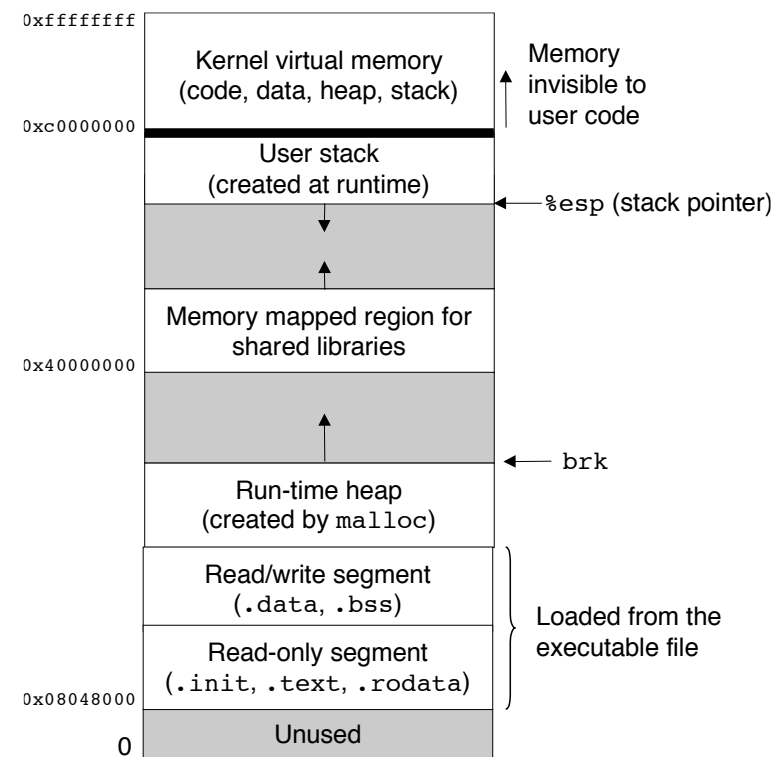
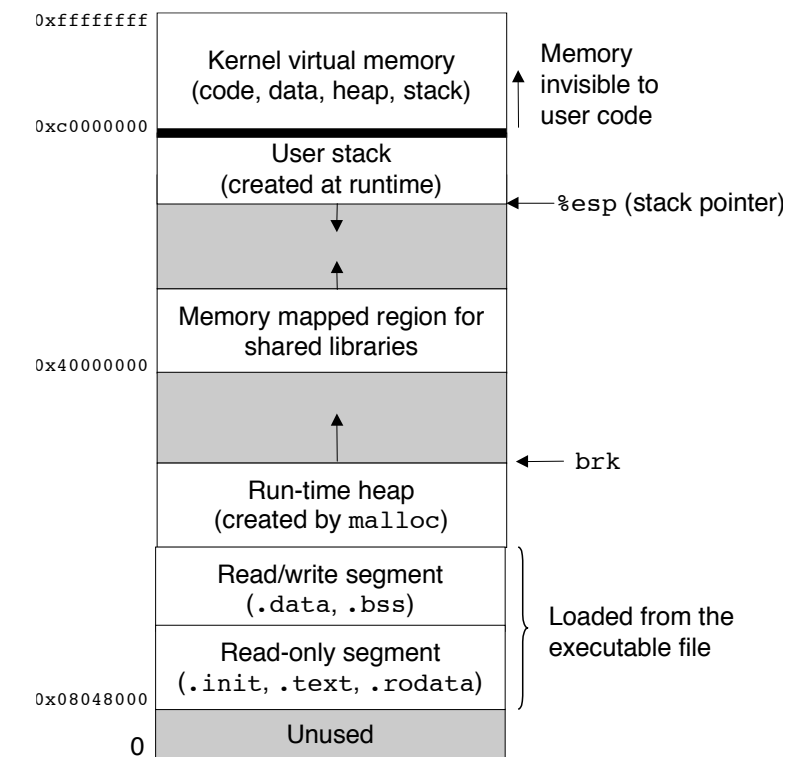


*address space consistency* → provide a uniform “view” of memory to each process



*memory protection* → prevent processes from directly accessing each other's address space



 $P_0$  $P_1$  $P_2$ 

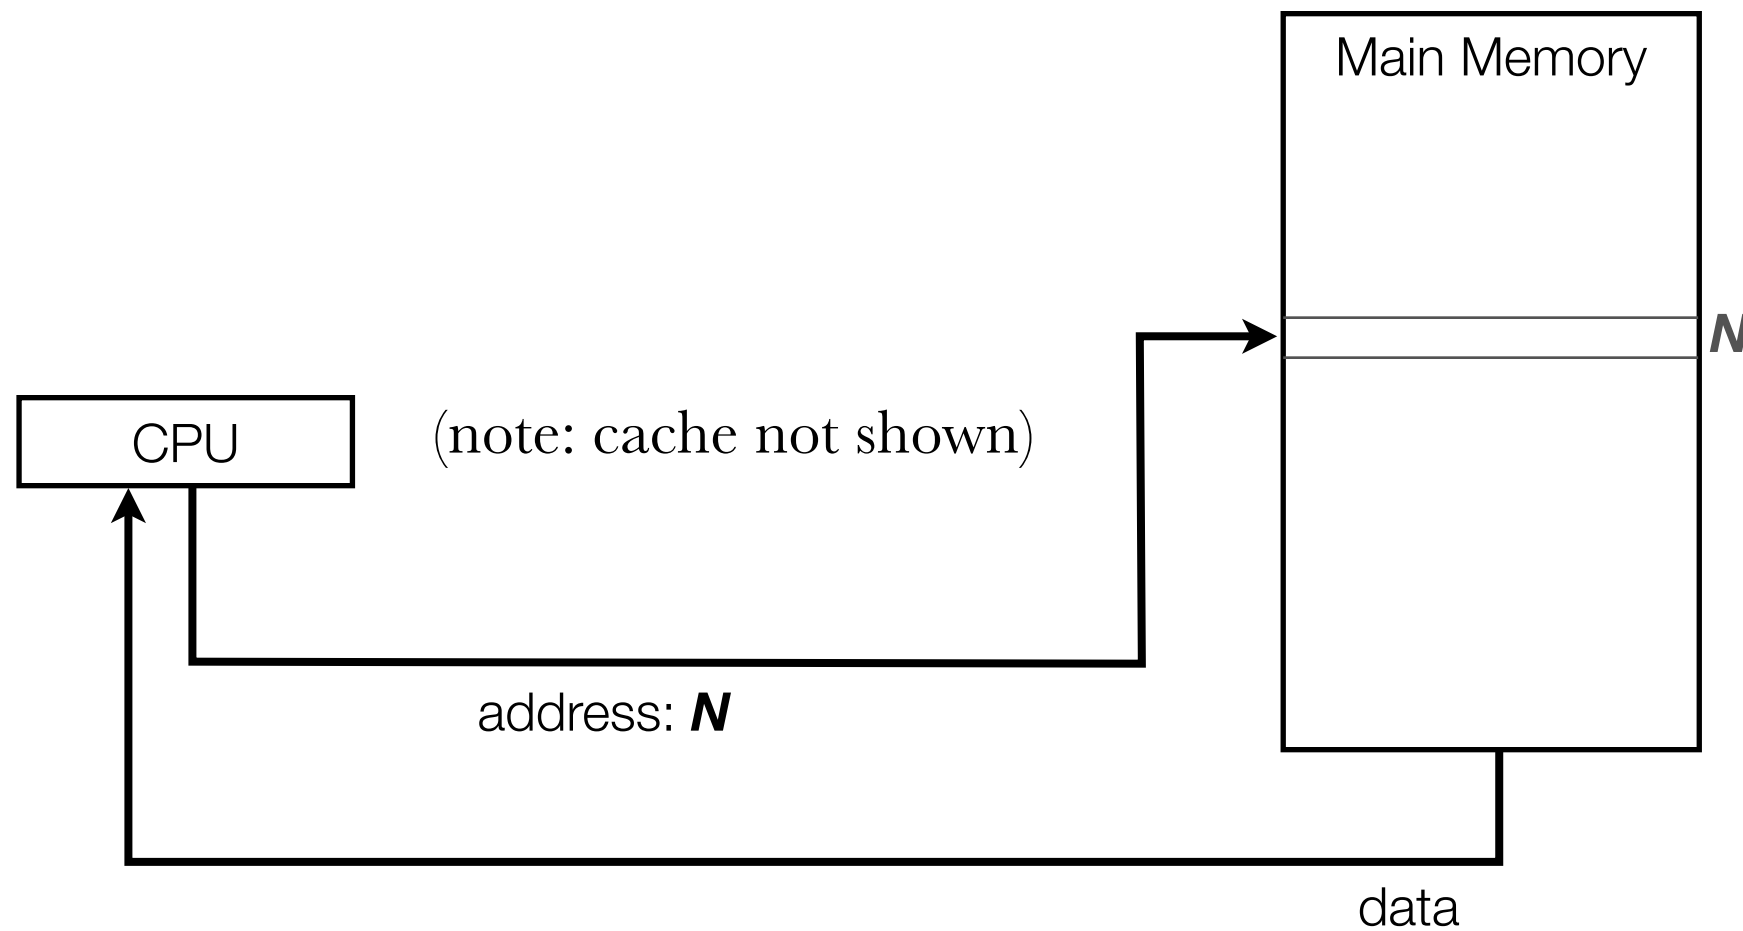
*memory protection* → prevent processes from directly accessing each other's address space

i.e., every process should be provided with a managed, *virtualized* address space



“memory addresses”: what are they, really?





“physical” address: (byte) index into DRAM



```
int glob = 0xDEADBEEE;
```

```
main() {
    fork();
    glob += 1;
}
```

**(gdb) set detach-on-fork off**

**(gdb) break main**

Breakpoint 1 at 0x400508: file memtest.c, line 7.

**(gdb) run**

Breakpoint 1, main () at memtest.c:7

```
7      fork();
```

**(gdb) next**

[New process 7450]

```
8      glob += 1;
```

**(gdb) print &glob**

\$1 = (int \*) 0x6008d4

**(gdb) next**

```
9      }
```

**(gdb) print /x glob**

\$2 = 0xdeadbeef

**(gdb) inferior 2**

[Switching to inferior 2 [process 7450]]

#0 0x000000310acac49d in \_\_libc\_fork ()

```
131     pid = ARCH_FORK ();
```

**(gdb) finish**

Run till exit from #0 in \_\_libc\_fork ()

```
8      glob += 1;
```

**(gdb) print /x glob**

\$4 = 0xdeadbeee

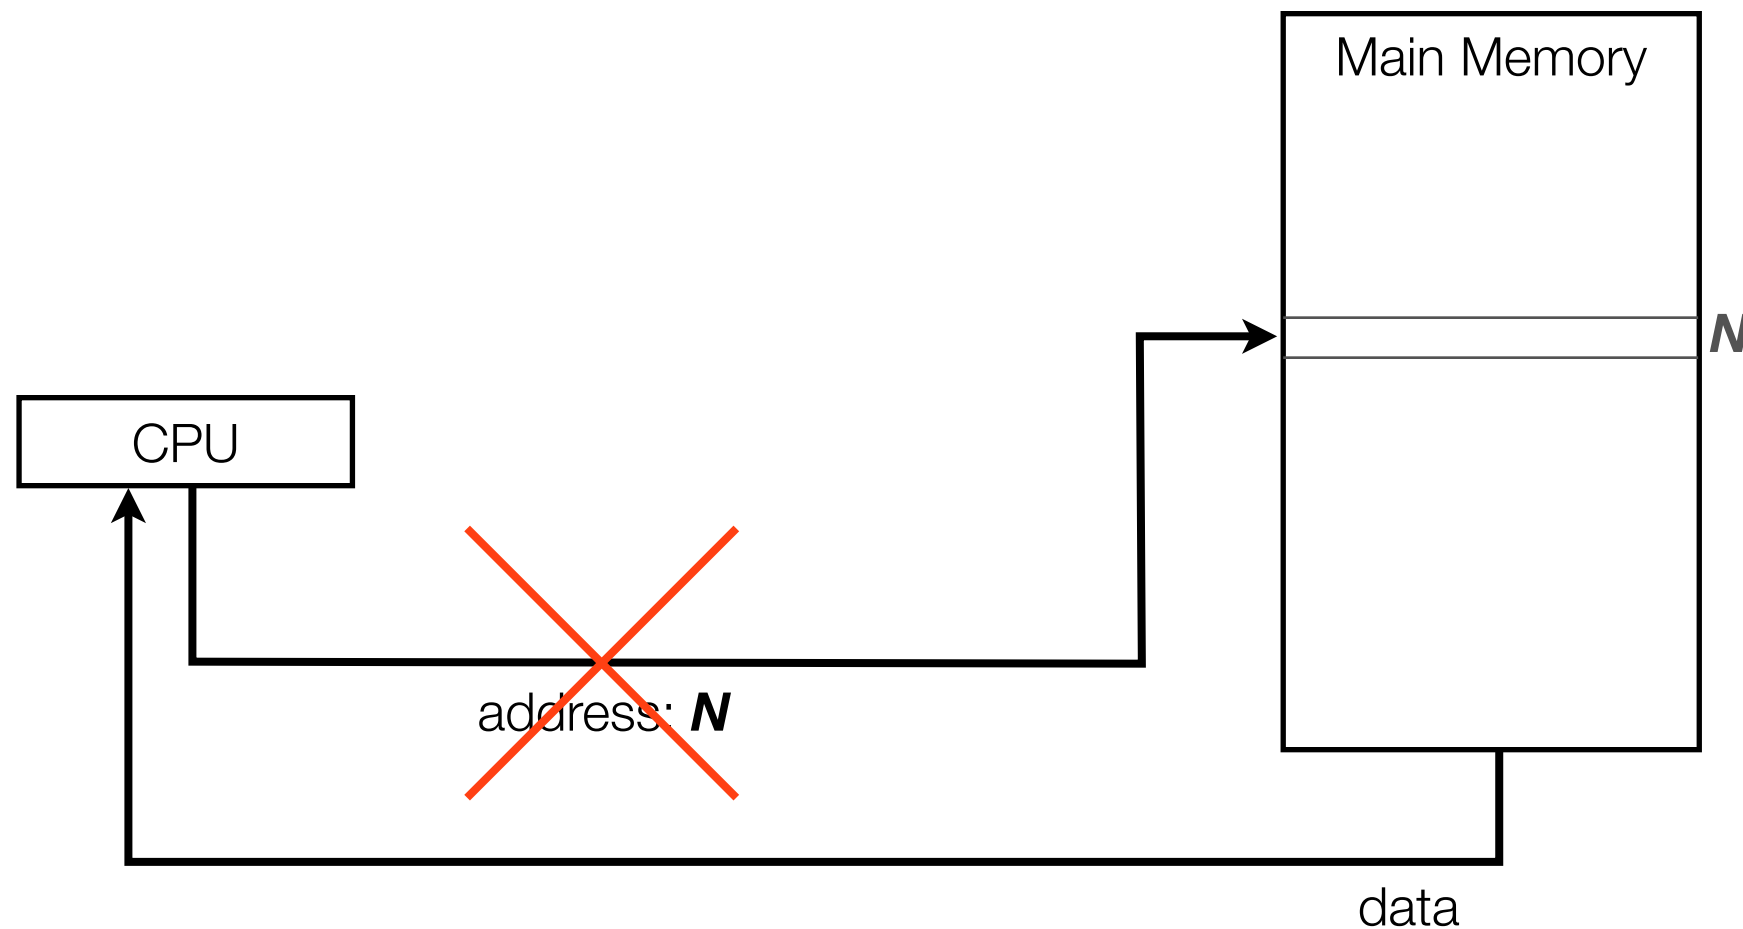
**(gdb) print &glob**

\$5 = (int \*) 0x6008d4

*parent*

*child*



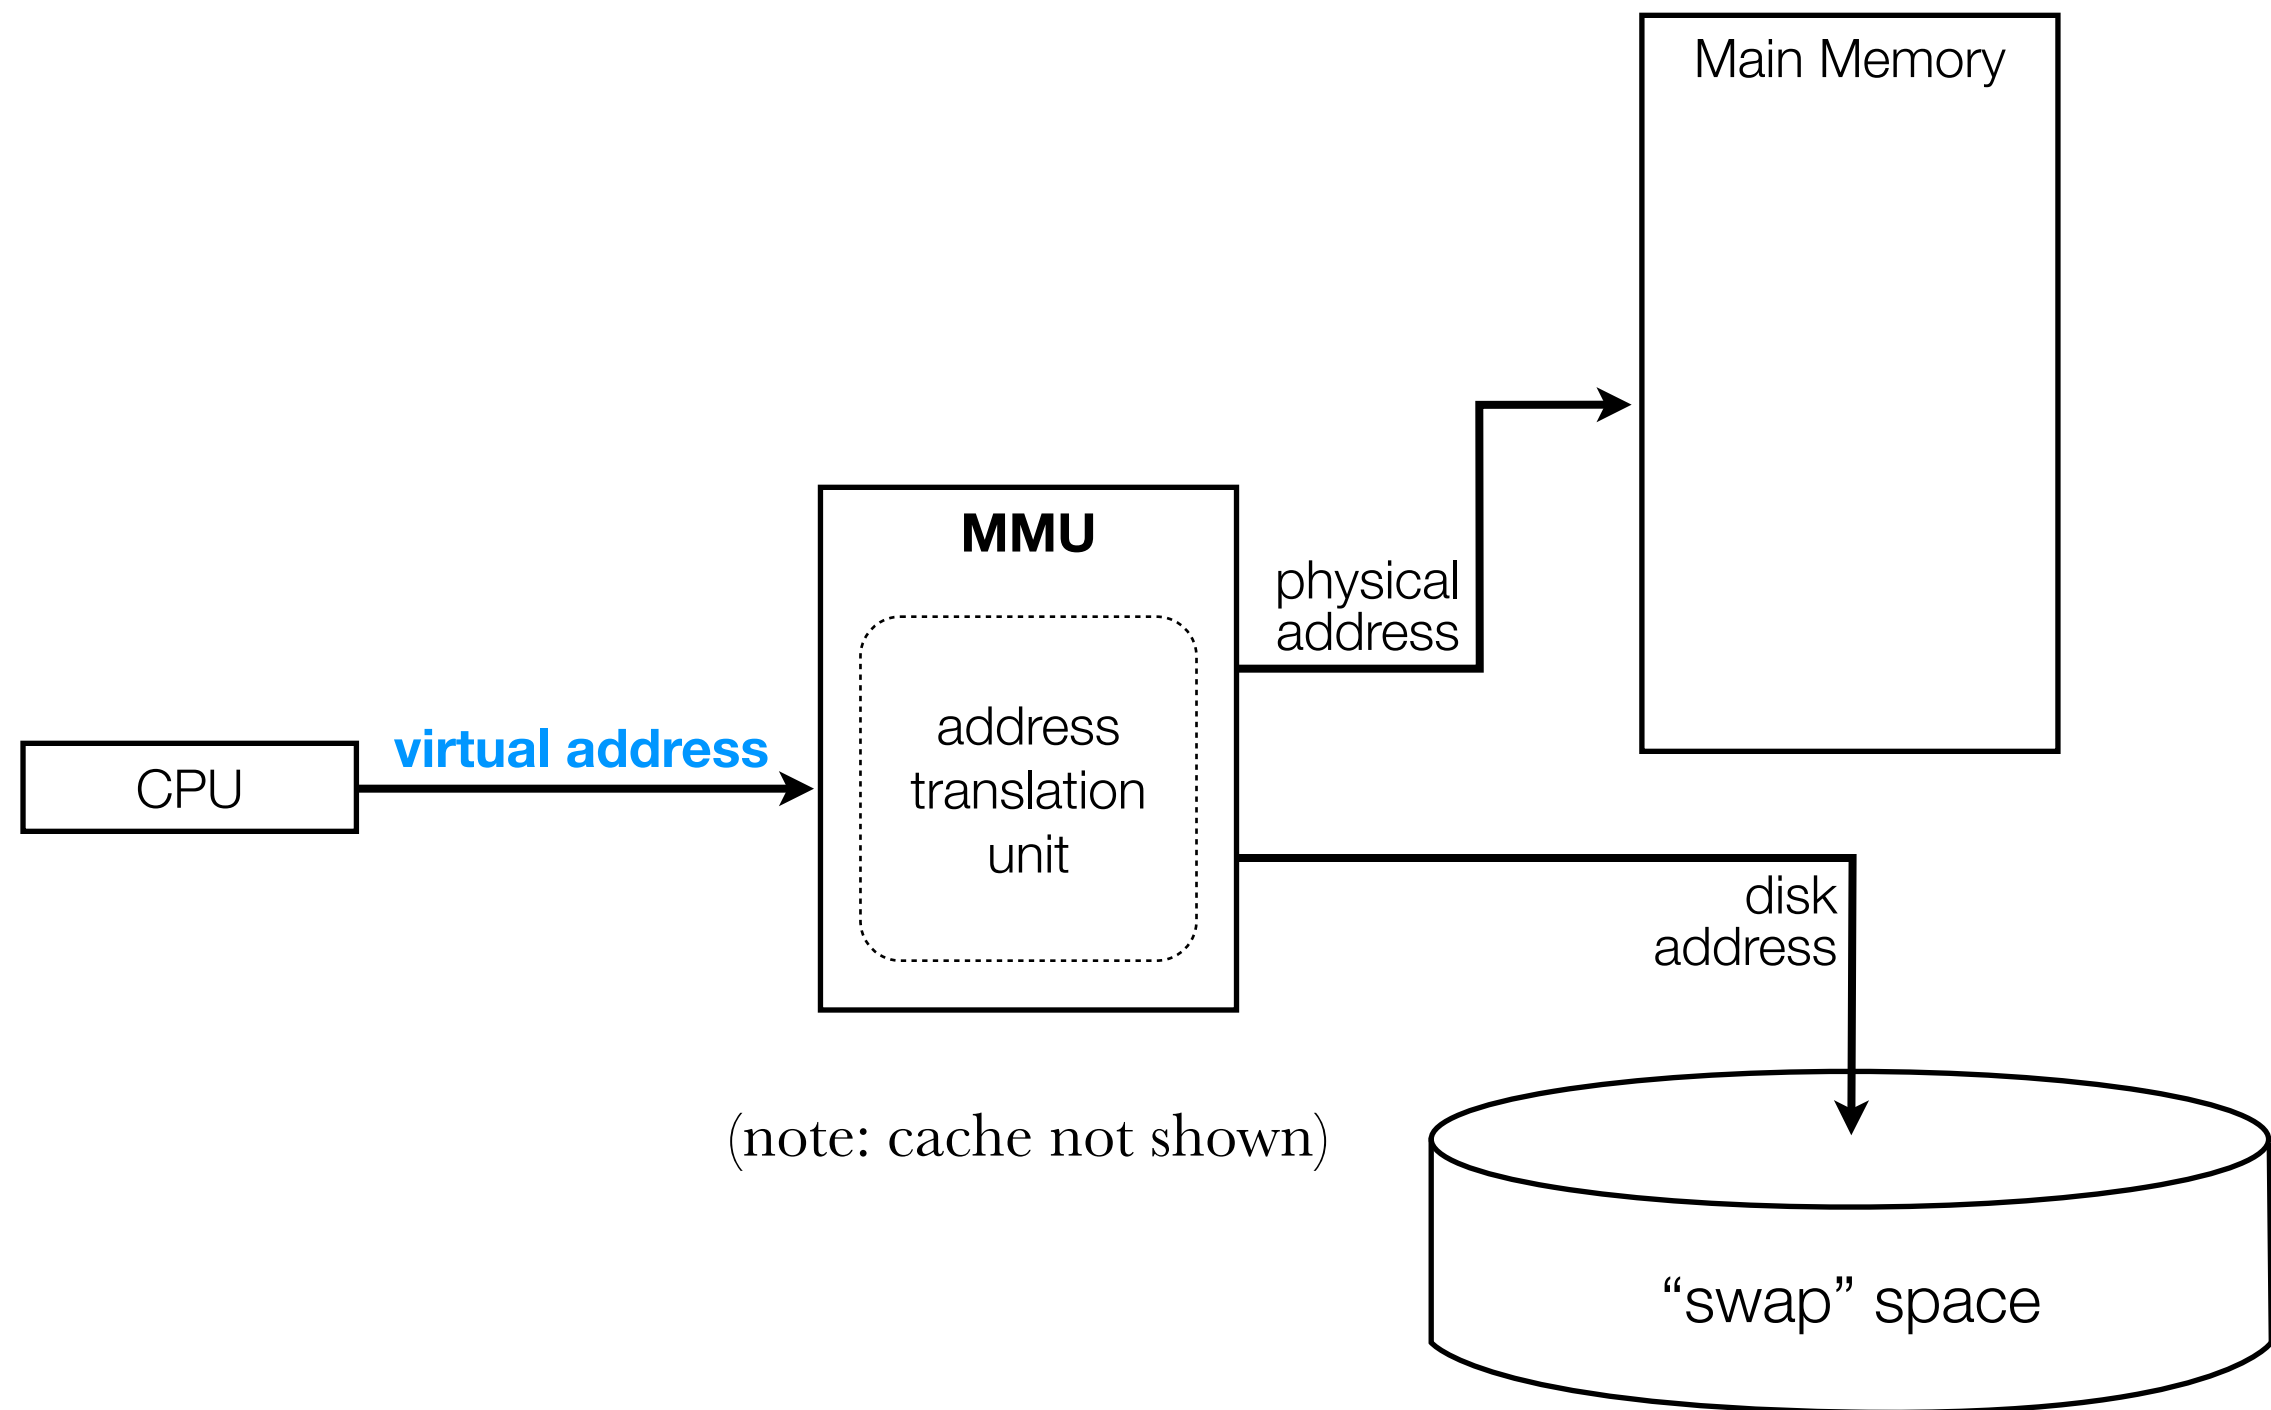


instructions executed by the CPU *do not* refer directly to *physical* addresses!

processes reference *virtual* addresses,  
the CPU relays virtual address requests to  
the *memory management unit* (MMU),  
which are *translated* to physical addresses







essential problem: translate request for a  
virtual address  $\rightarrow$  physical address

... this must be **FAST**, as *every* memory  
access from the CPU must be translated



both hardware/software are involved:


- MMU (hw) handles simple and fast operations (e.g., table lookups)
- Kernel (sw) handles complex tasks (e.g., eviction policy)




# § Virtual Memory Implementations



# 1. simple relocation

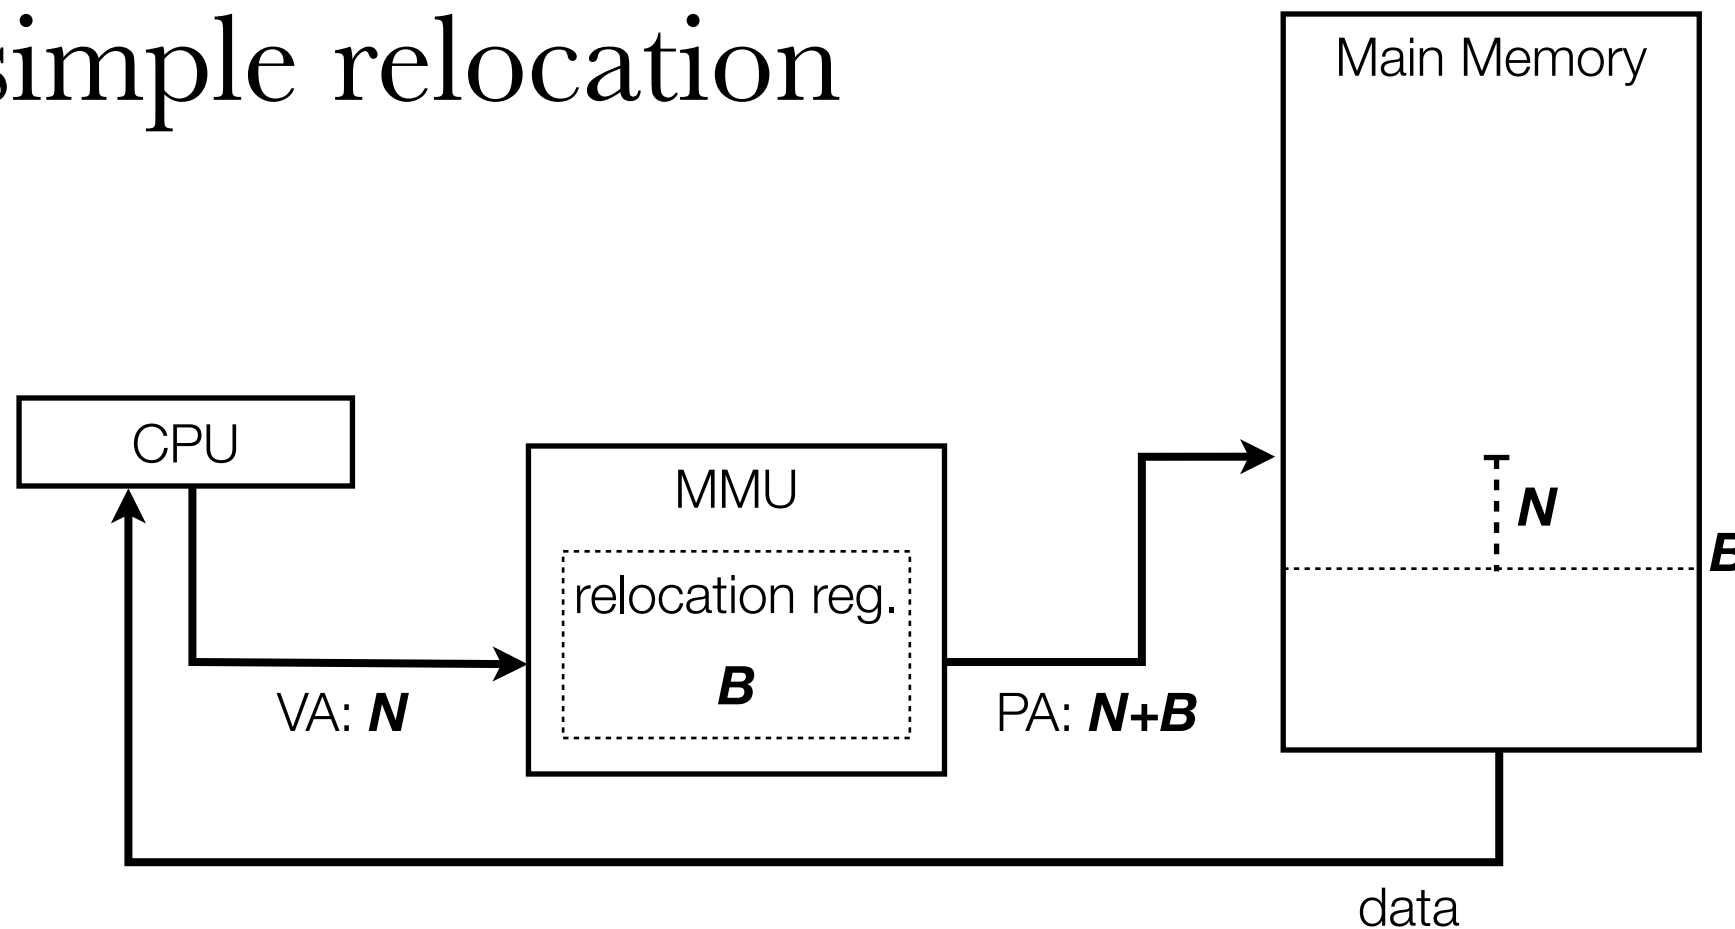


# 1. simple relocation



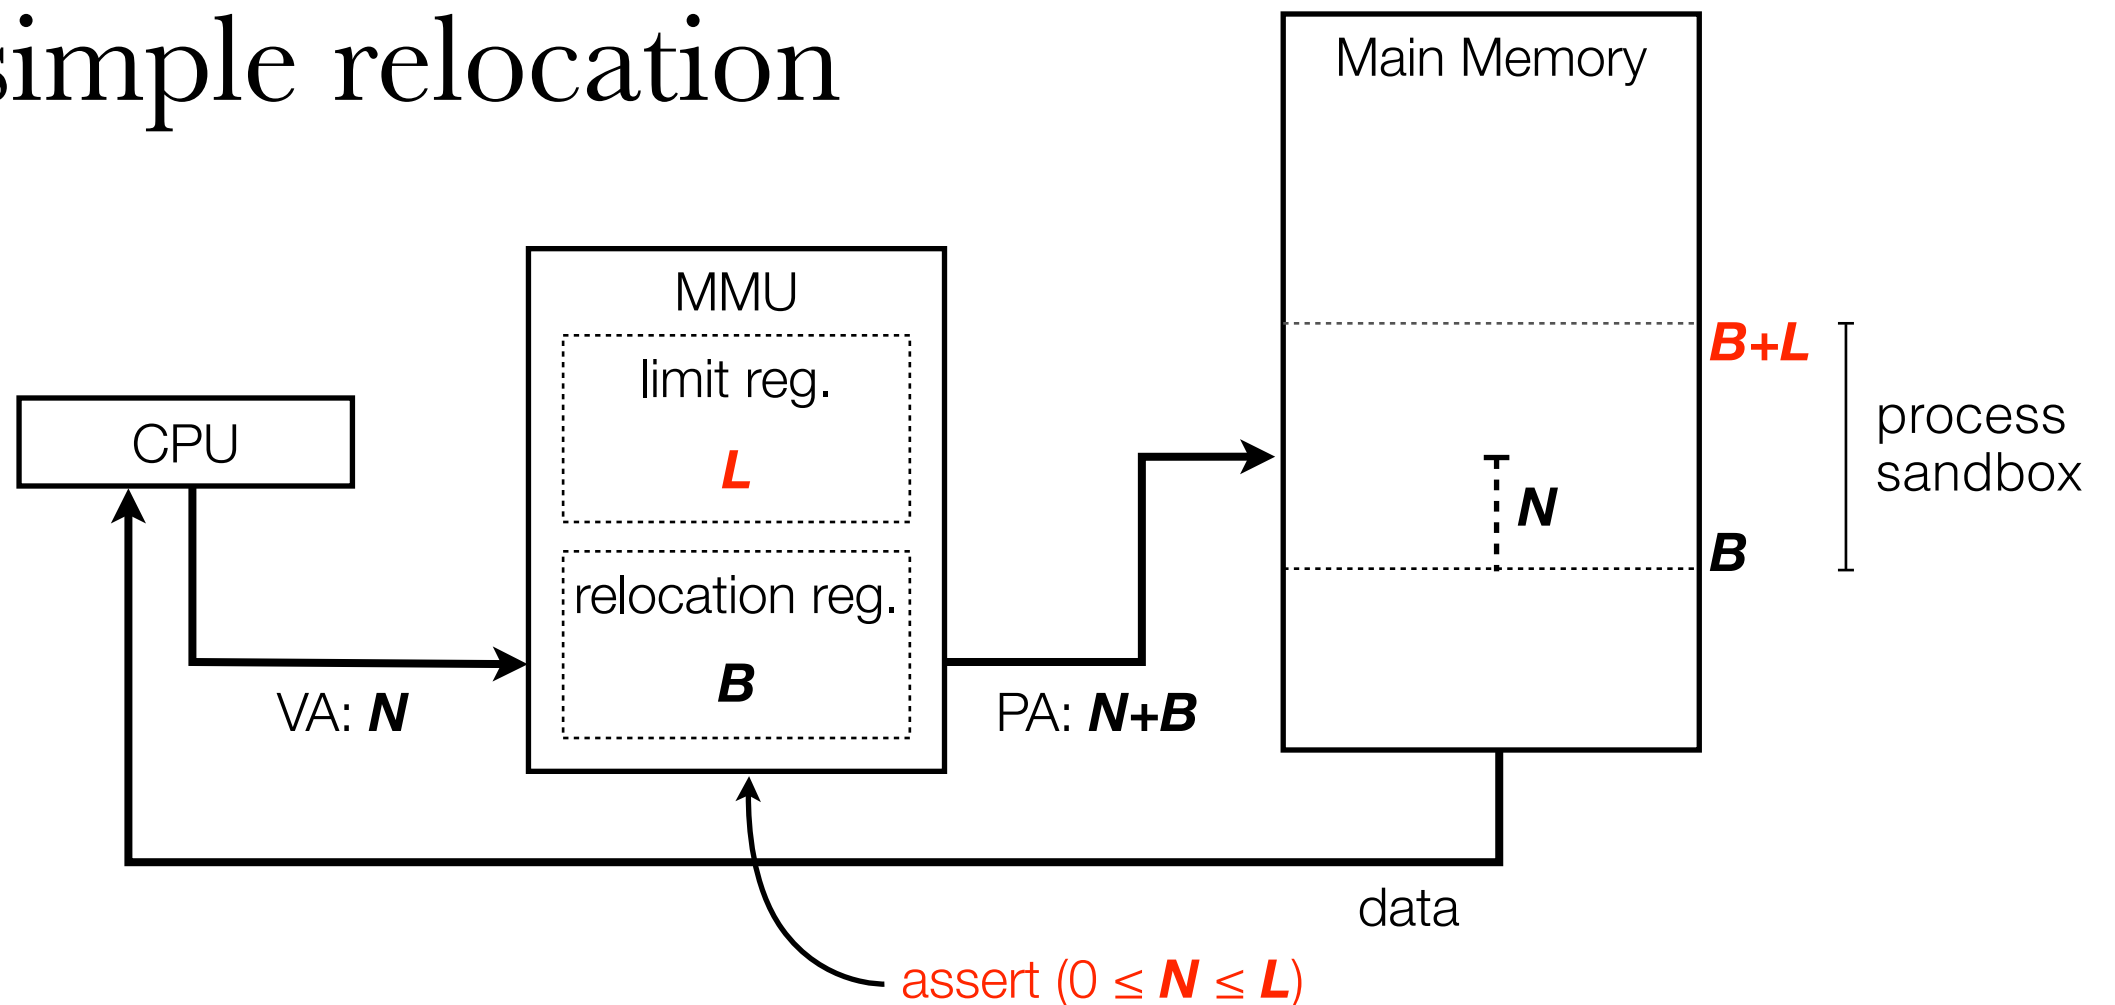
- per-process relocation address is loaded by kernel on every context switch

# 1. simple relocation



- problem: processes may easily overextend their bounds and trample on each other


# 1. simple relocation



- incorporate a *limit* register to provide memory protection



# 1. simple relocation

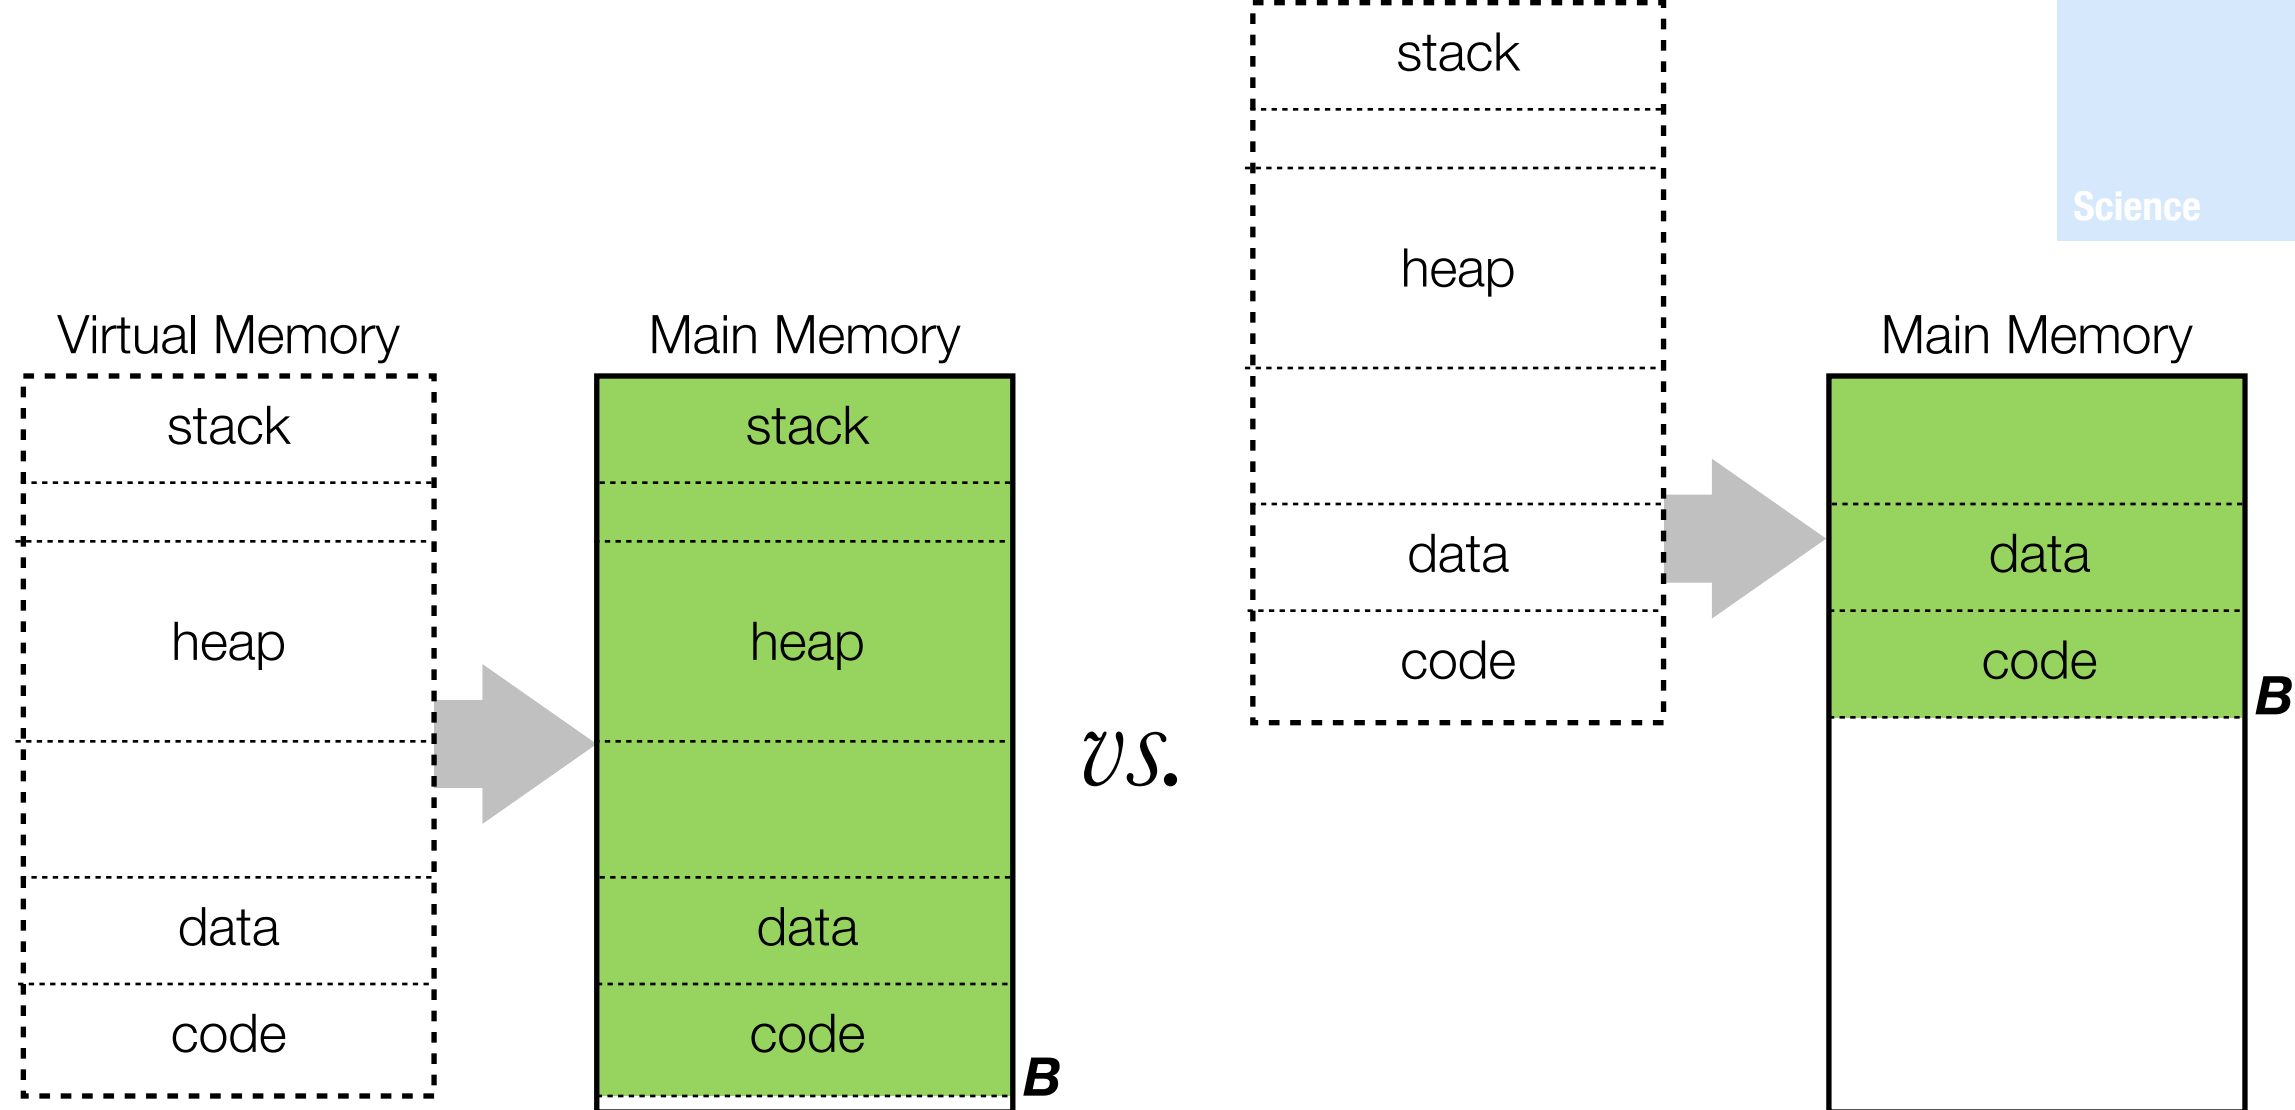


- assertion failure triggers a fault, which summons kernel (which signals process)

pros:


- simple & fast!
- provides protection





but: available memory for mapping  
depends on value of base address

i.e., address spaces are *not consistent*!



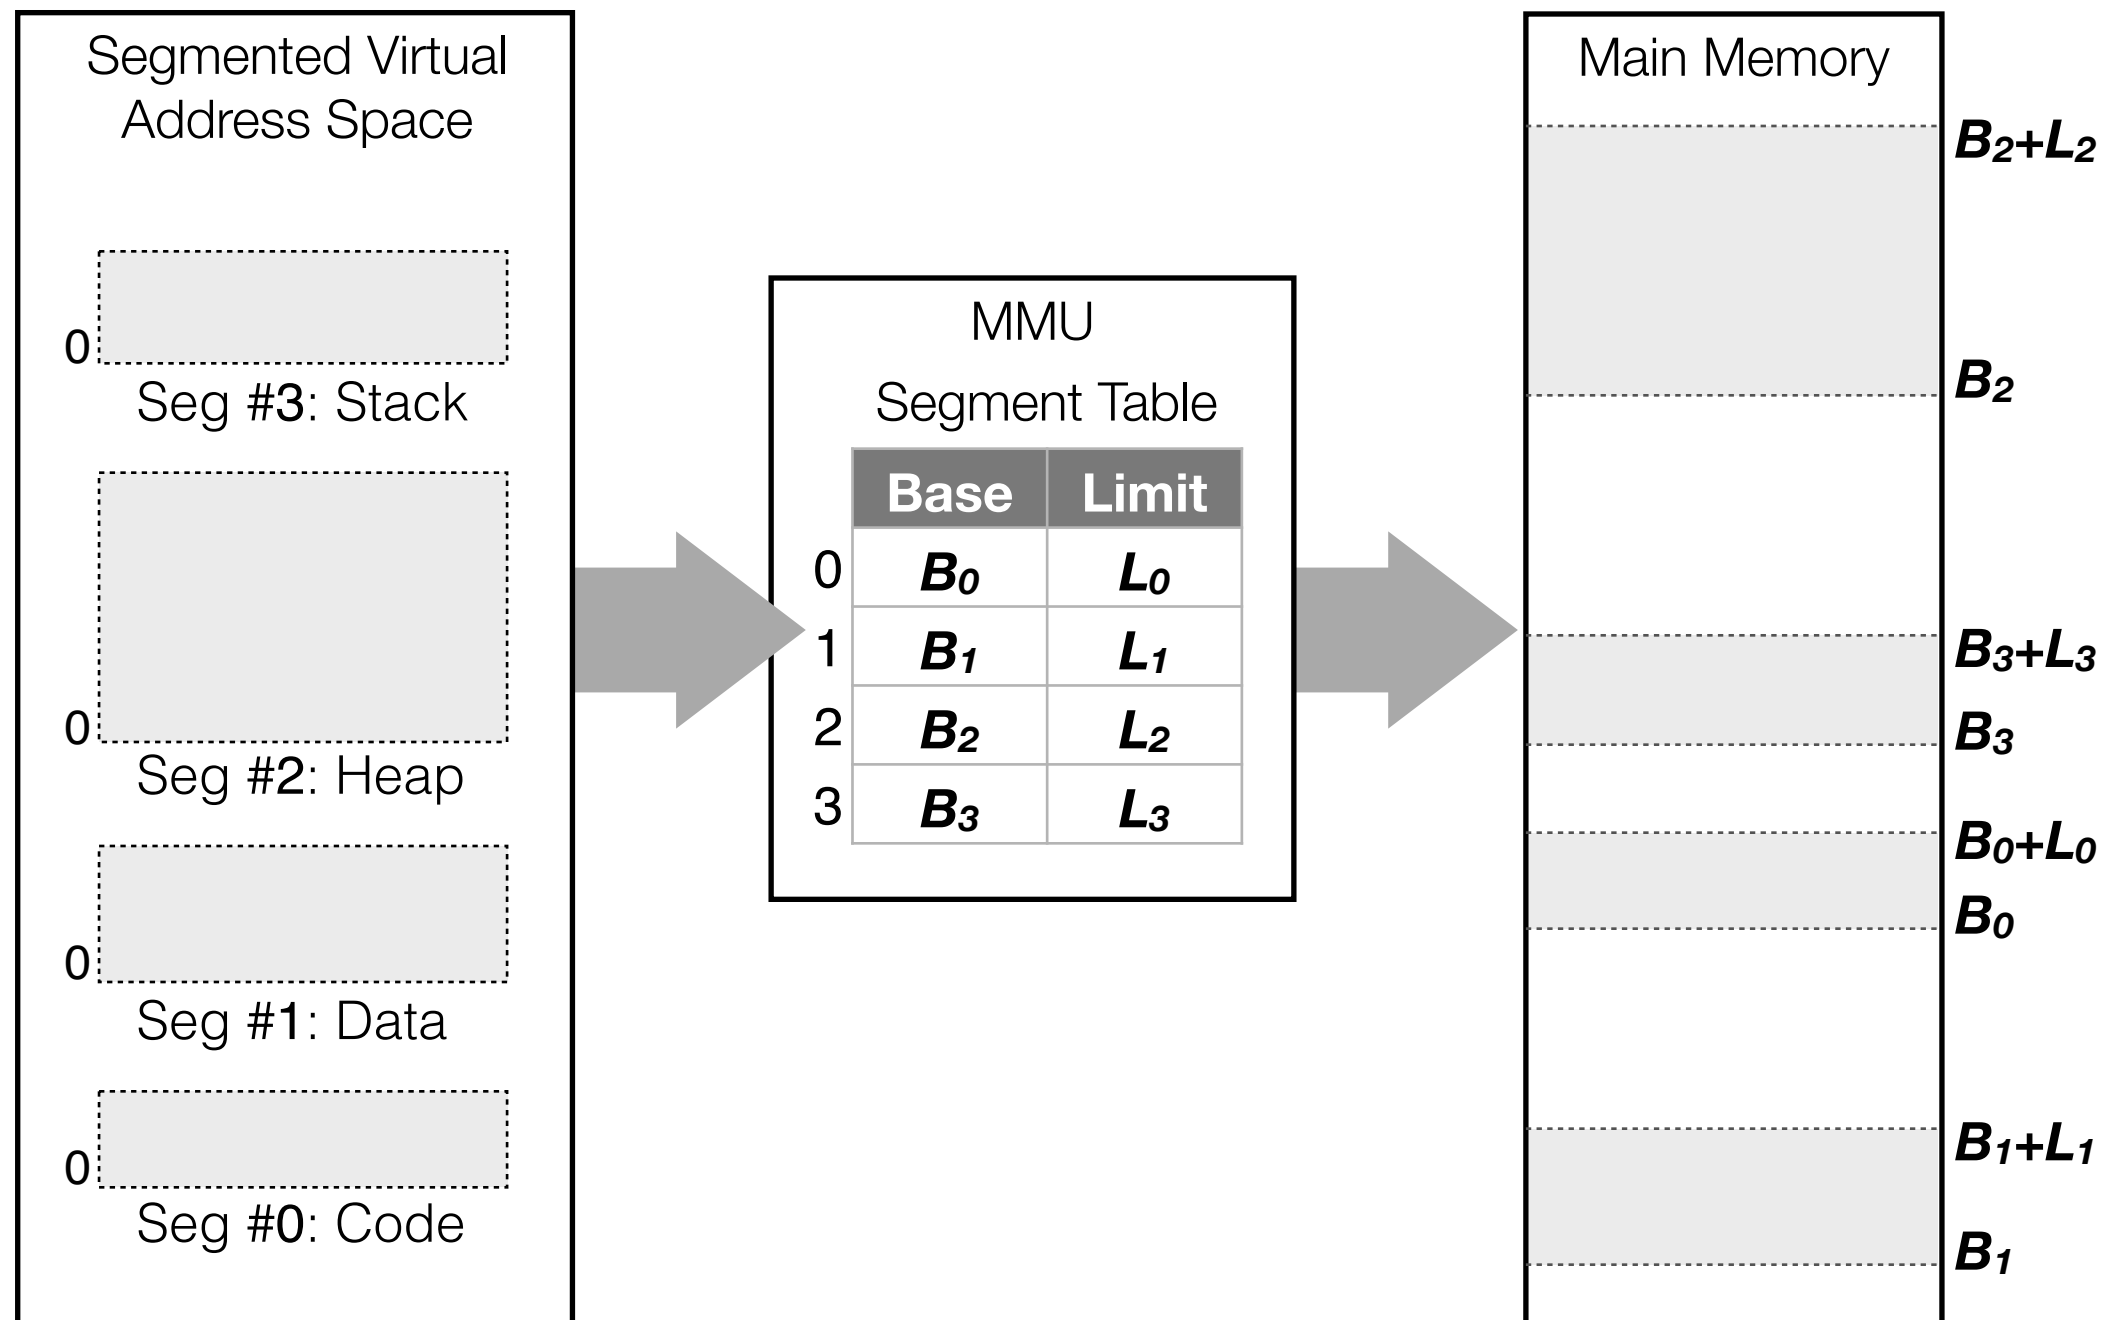
also: all of a process *below the address limit*  
must be loaded in memory

i.e., memory may be *vastly under-utilized*

## 2. segmentation

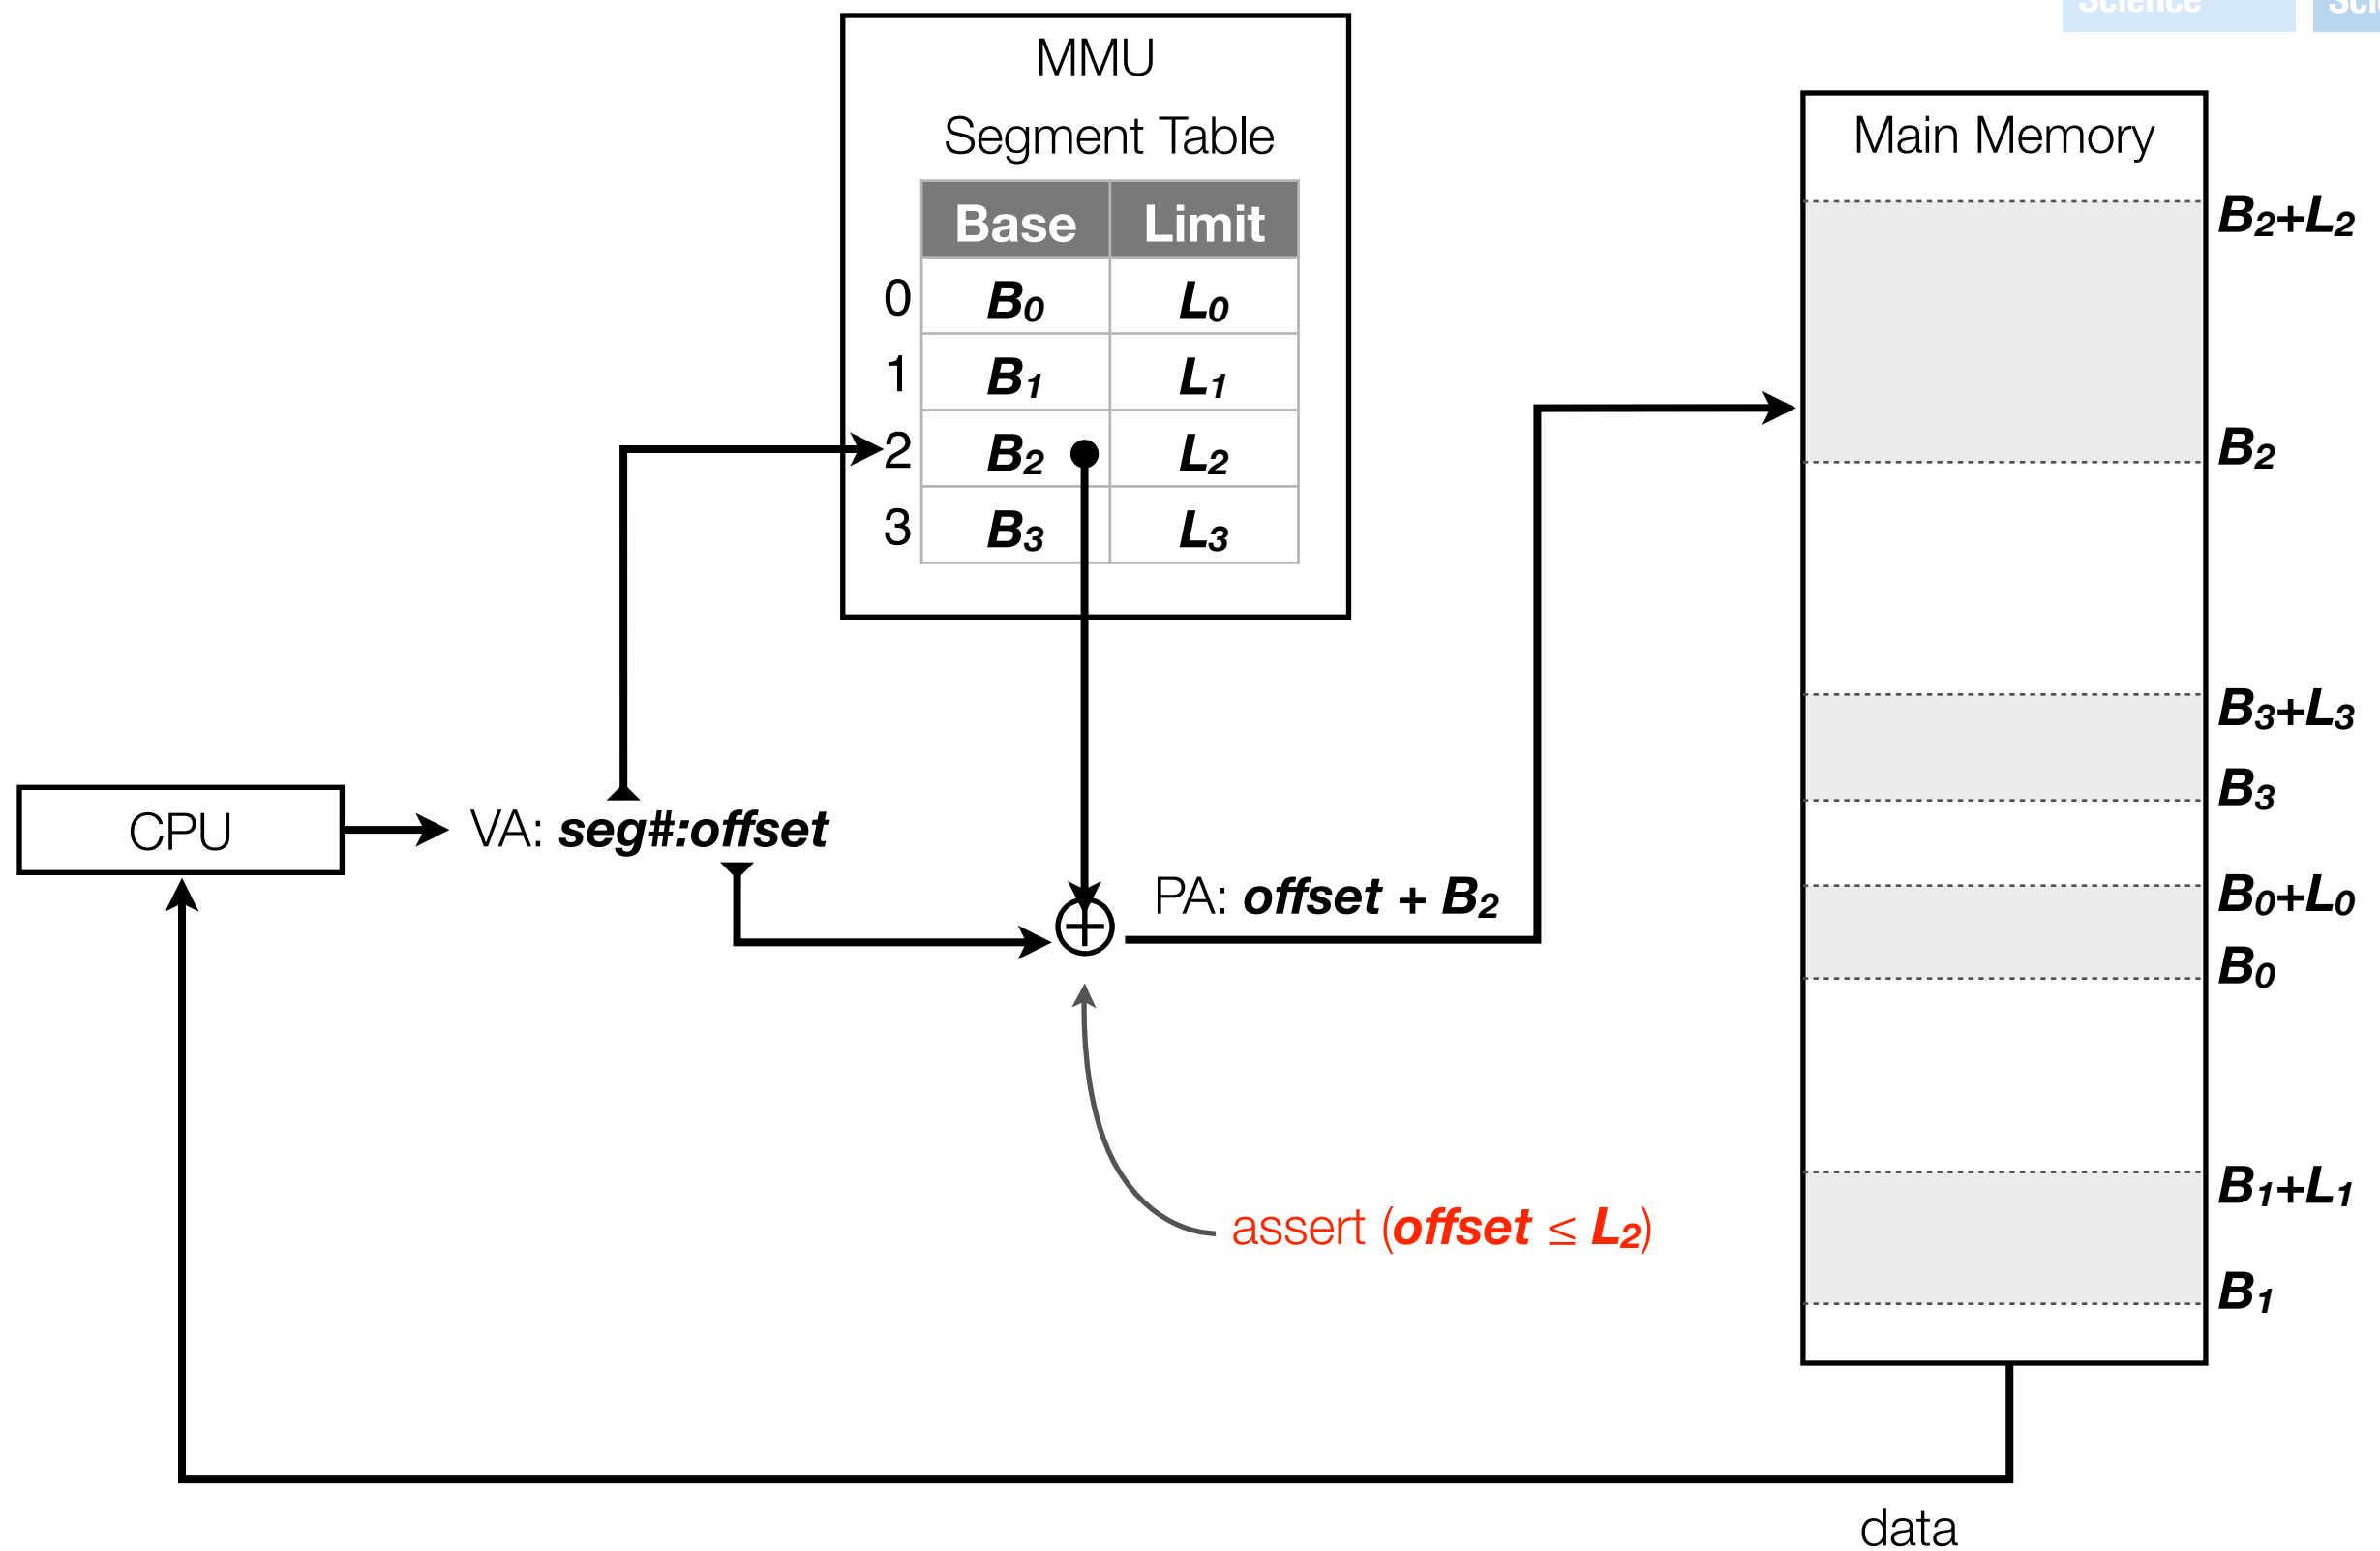
- partition virtual address space into *multiple logical segments*
- individually map them onto physical memory with relocation registers





virtual address has form *seg#.offset*





Segment Table

	Base	Limit
0	$B_0$	$L_0$
1	$B_1$	$L_1$
2	$B_2$	$L_2$
3	$B_3$	$L_3$

- implemented as MMU registers
- part of kernel-maintained, per-process metadata (aka “process control block”)
- re-populated on each context switch






pros:

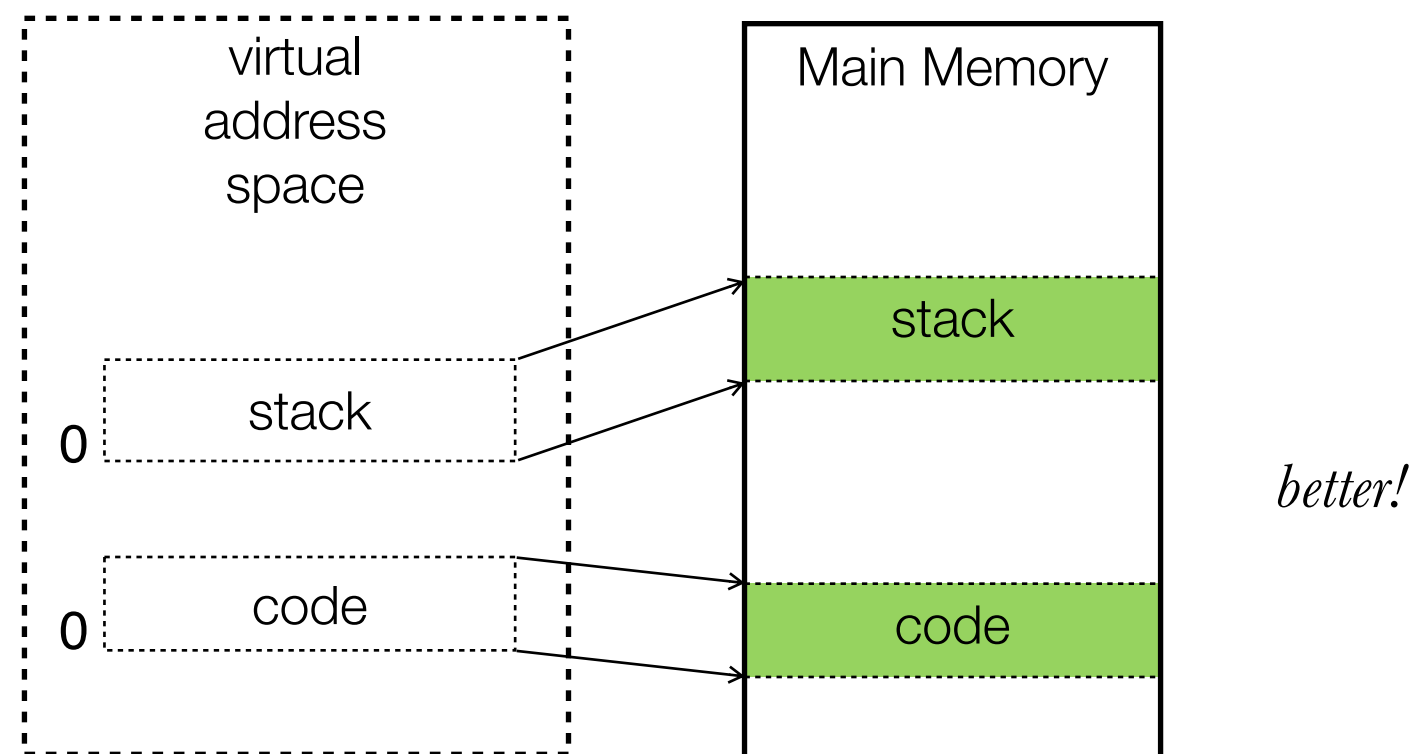
- still very *fast*
- translation = register access & addition
- memory *protection* via limits
- segmented addresses improve *consistency*




*simple  
relocation:*



*segmentation:*







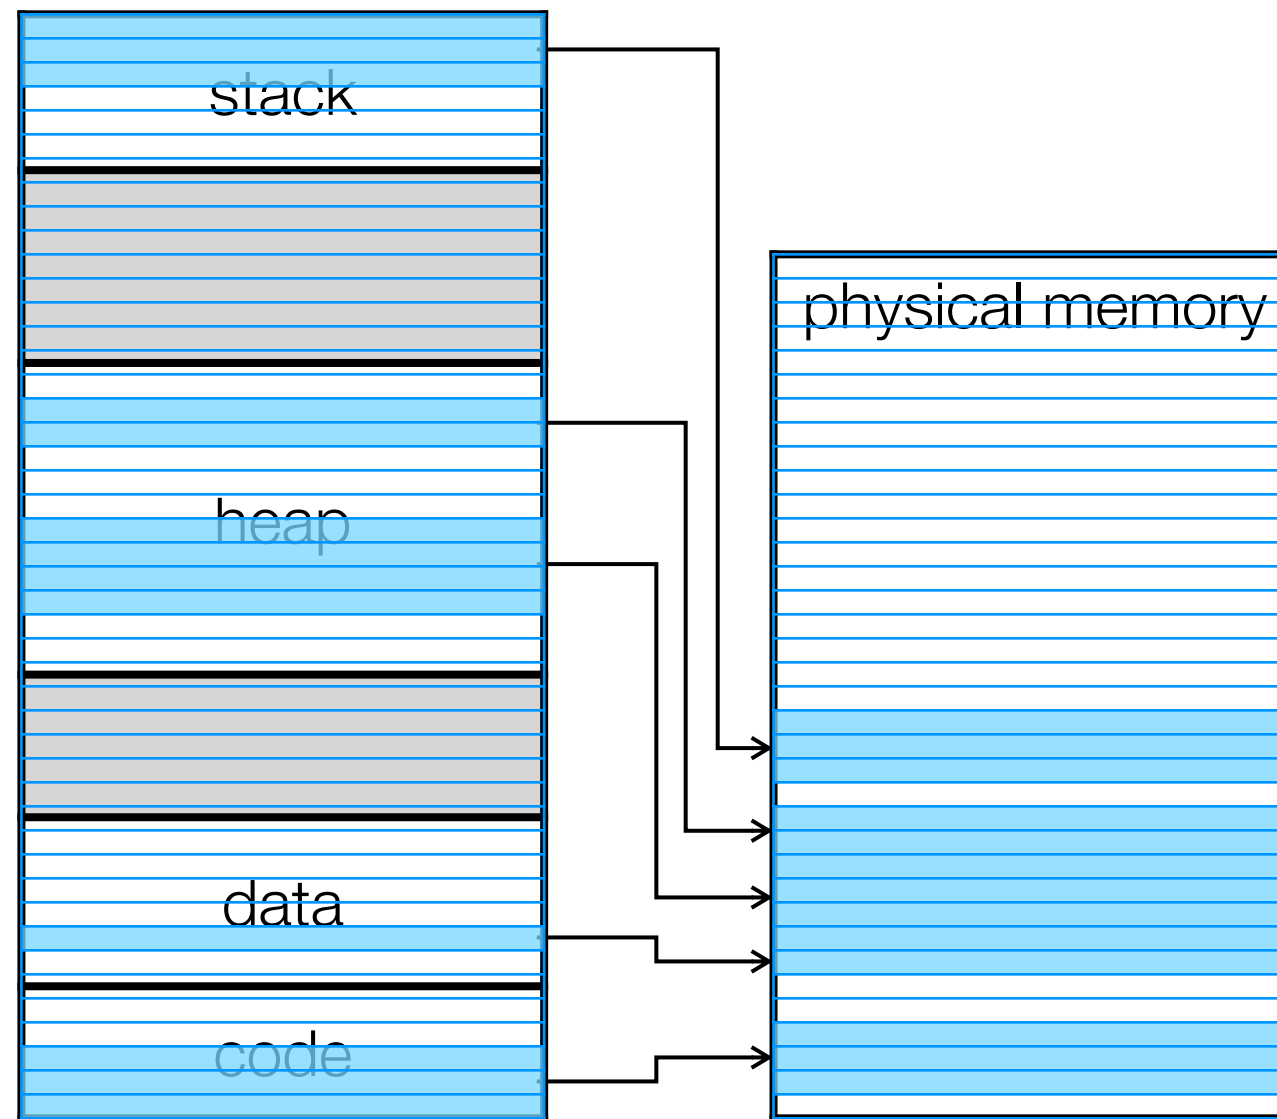
- variable segment sizes  $\rightarrow$  *memory fragmentation*
- fragmentation potentially *lowers utilization*
- can fix through compaction, but expensive!

### 3. paging

- partition virtual and physical address spaces into *uniformly sized* **pages**
- virtual pages map onto physical pages







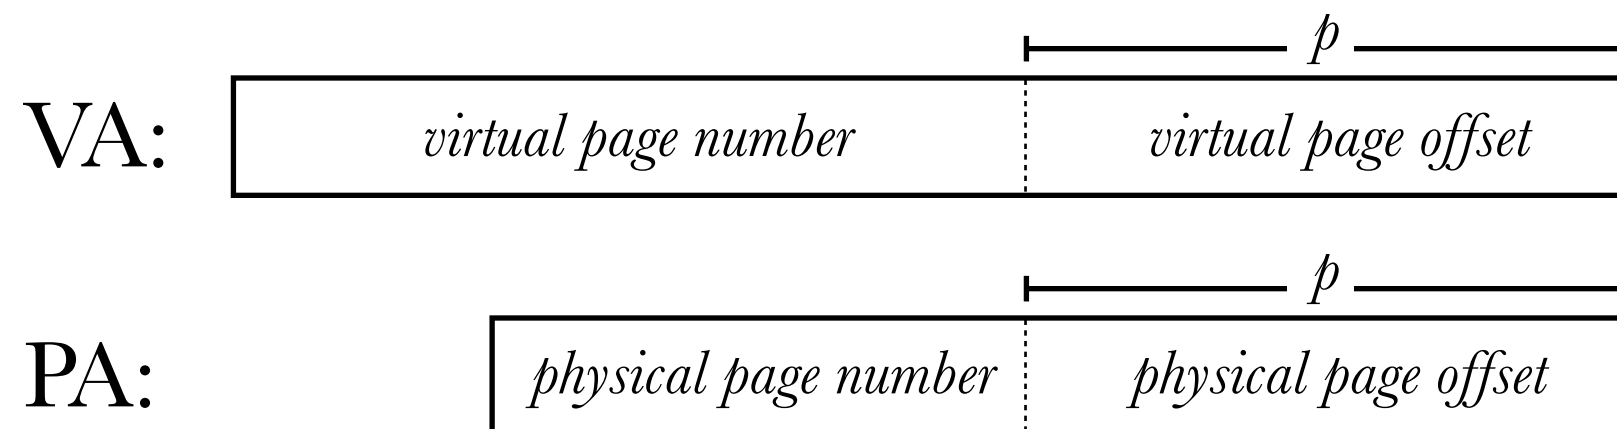
- minimum mapping granularity = page
- not all of a given segment need be mapped

modified mapping problem:


- a virtual address is broken down into *virtual page number & page offset*
- determine which physical page (if any) a given virtual page is loaded into
- if physical page is found, use page offset to access data




Given page size =  $2^p$  bytes









page table entries (PTEs) typically contain additional metadata, e.g.:

- dirty (modified) bit
- access bits (shared or kernel-owned pages may be read-only or inaccessible)



e.g., 32-bit virtual address,  
4KB ( $2^{12}$ ) page size,  
4-byte PTE size;  
- size of page table?



e.g., 32-bit virtual address,  
4KB ( $2^{12}$ ) pages,  
4-byte PTEs;

- # pages =  $2^{32} \div 2^{12} = 2^{20} = 1\text{M}$

- page table size =  $1\text{M} \times 4 \text{ bytes} = \mathbf{4\text{MB}}$



4MB is much too large to fit in the MMU  
— insufficient registers and SRAM!

Page table resides in **main memory**



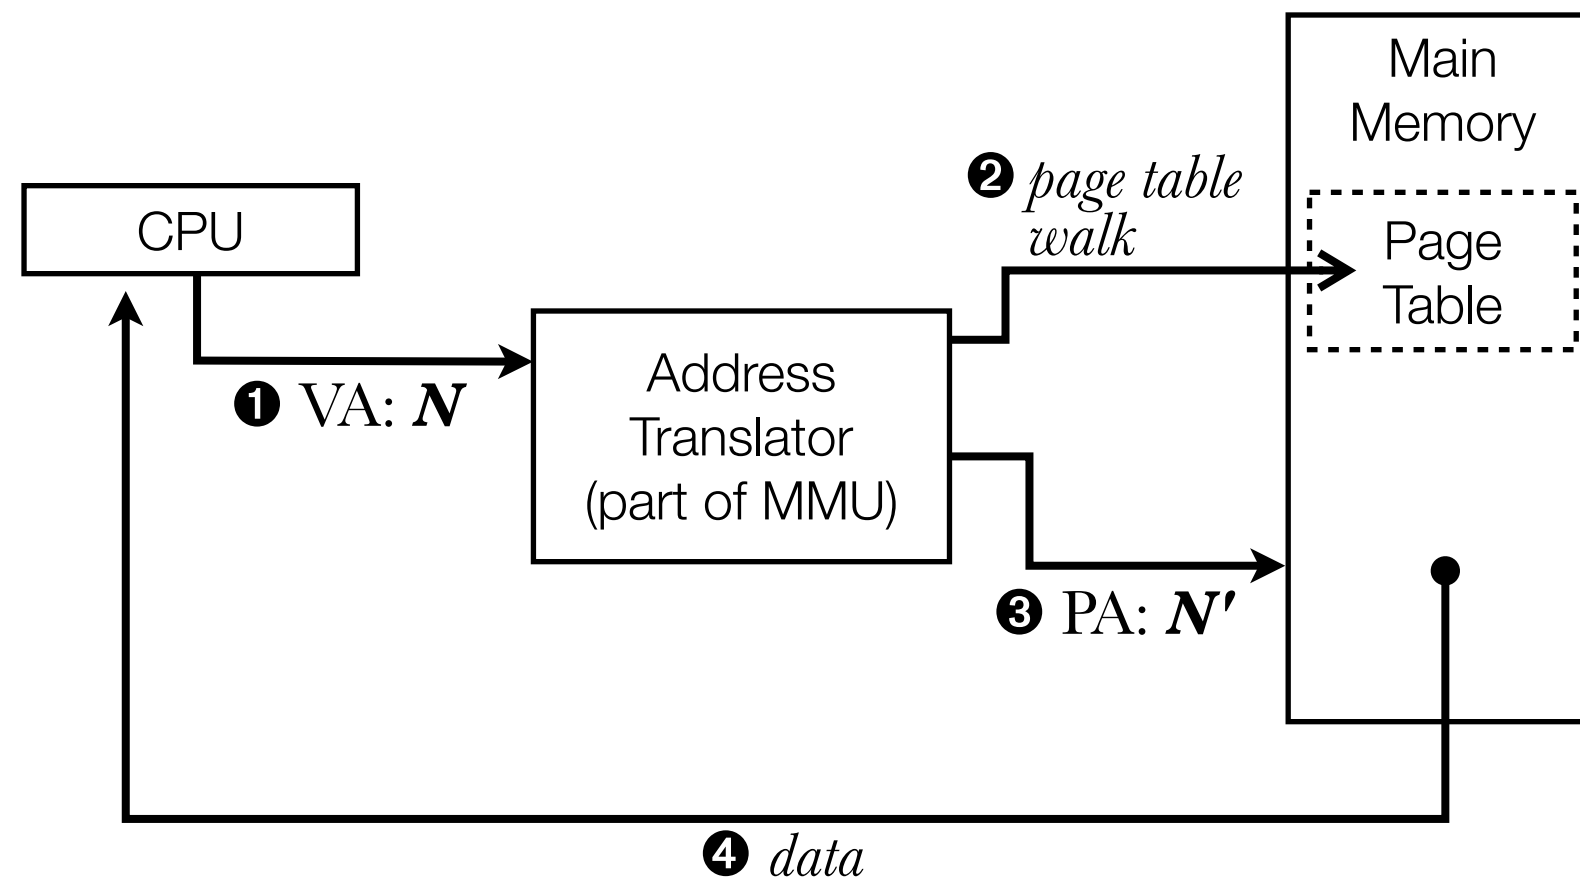
The translation process (aka *page table walk*) is performed by hardware (MMU).

The kernel must initially populate, then continue to manage a process's page table

The kernel also populates a *page table base register* on context switches

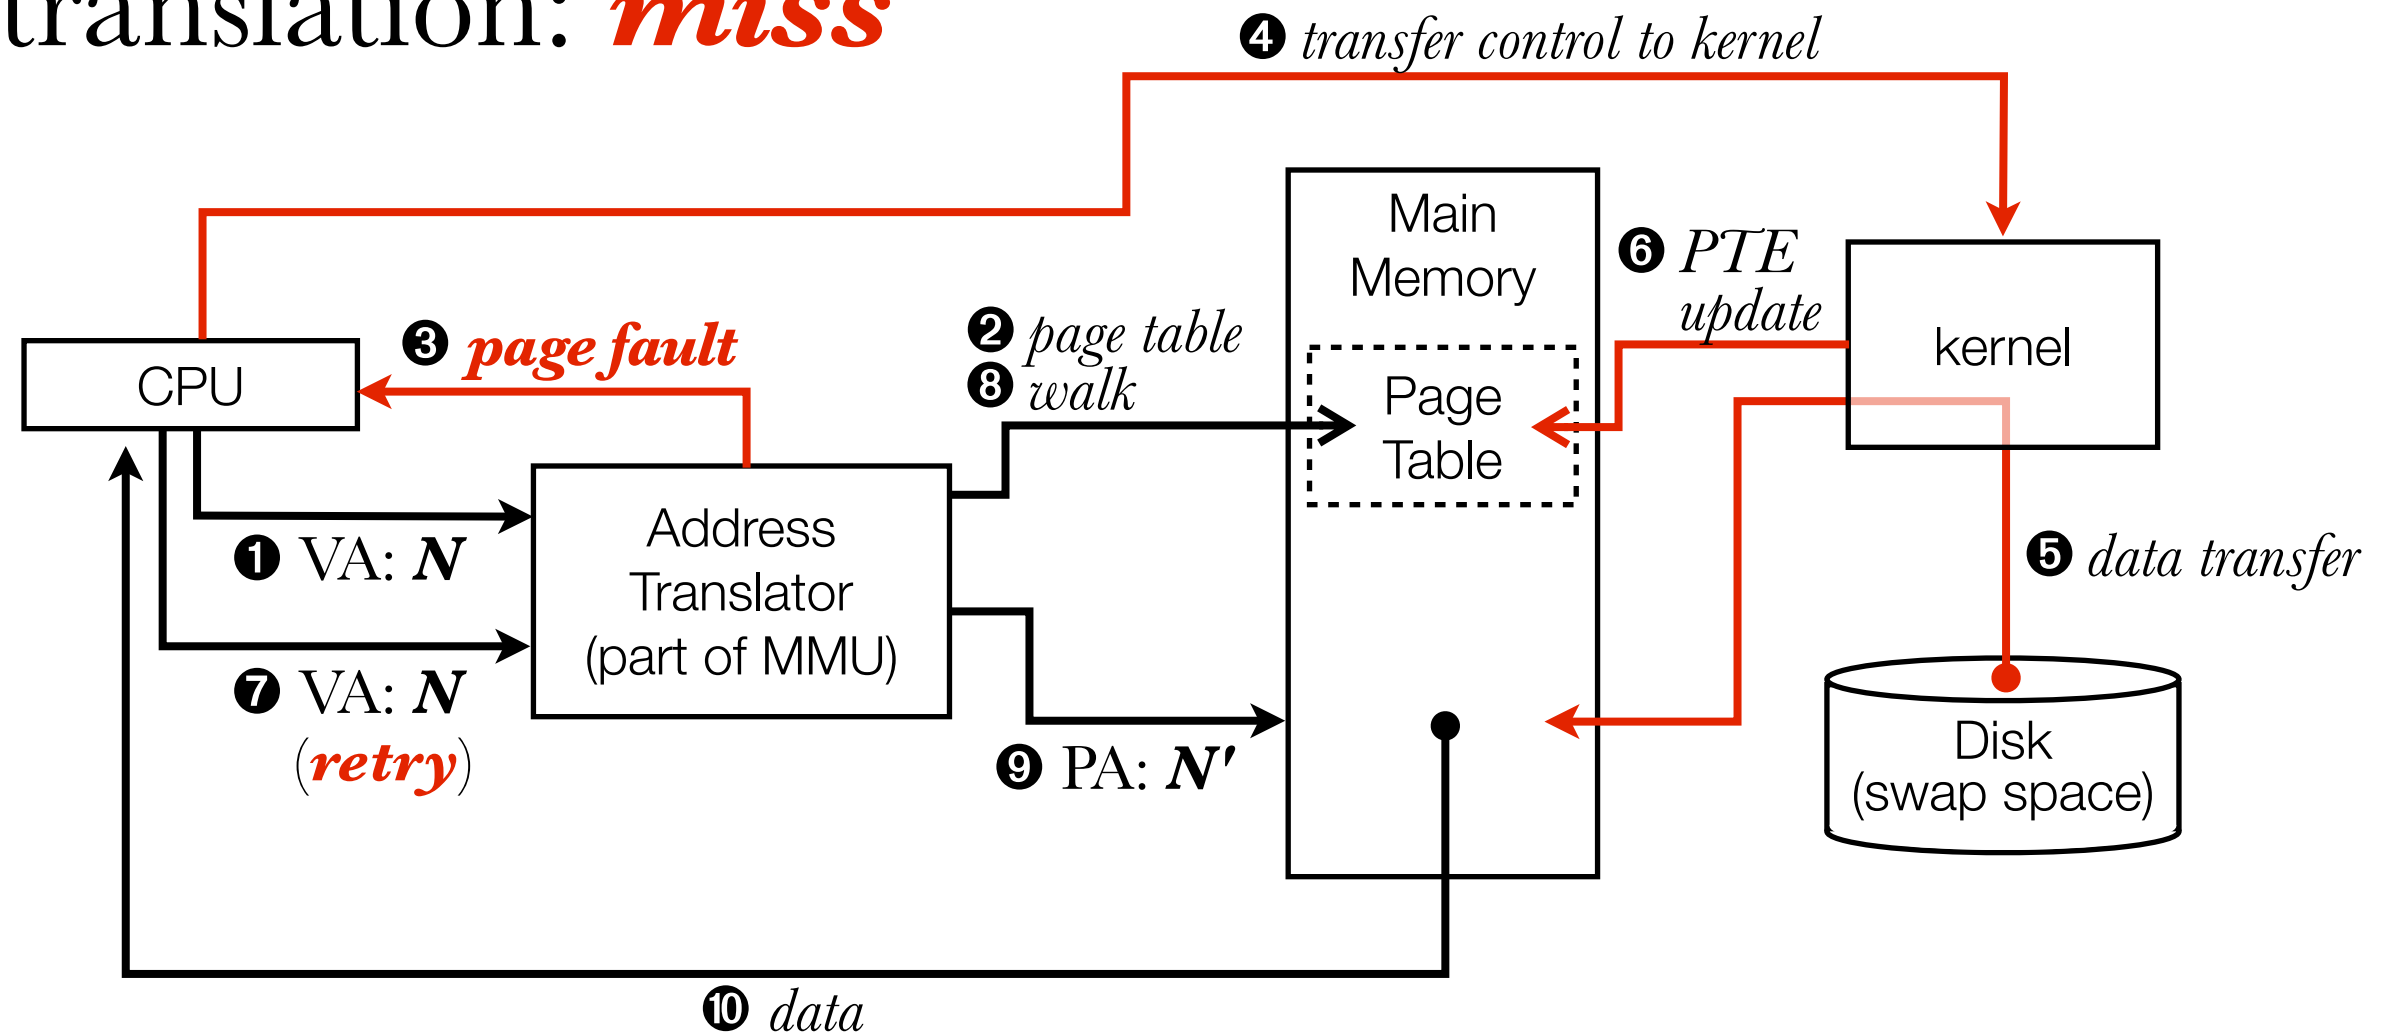


translation: *hit*





translation: *miss*



kernel decides where to place page, and what to evict (if memory is full)

- e.g., using LRU replacement policy



this system enables **on-demand paging**  
i.e., an active process need only be partly in  
memory (load rest from disk dynamically)



but if working set (of active processes)  
exceeds available memory, we may have  
**swap thrashing**



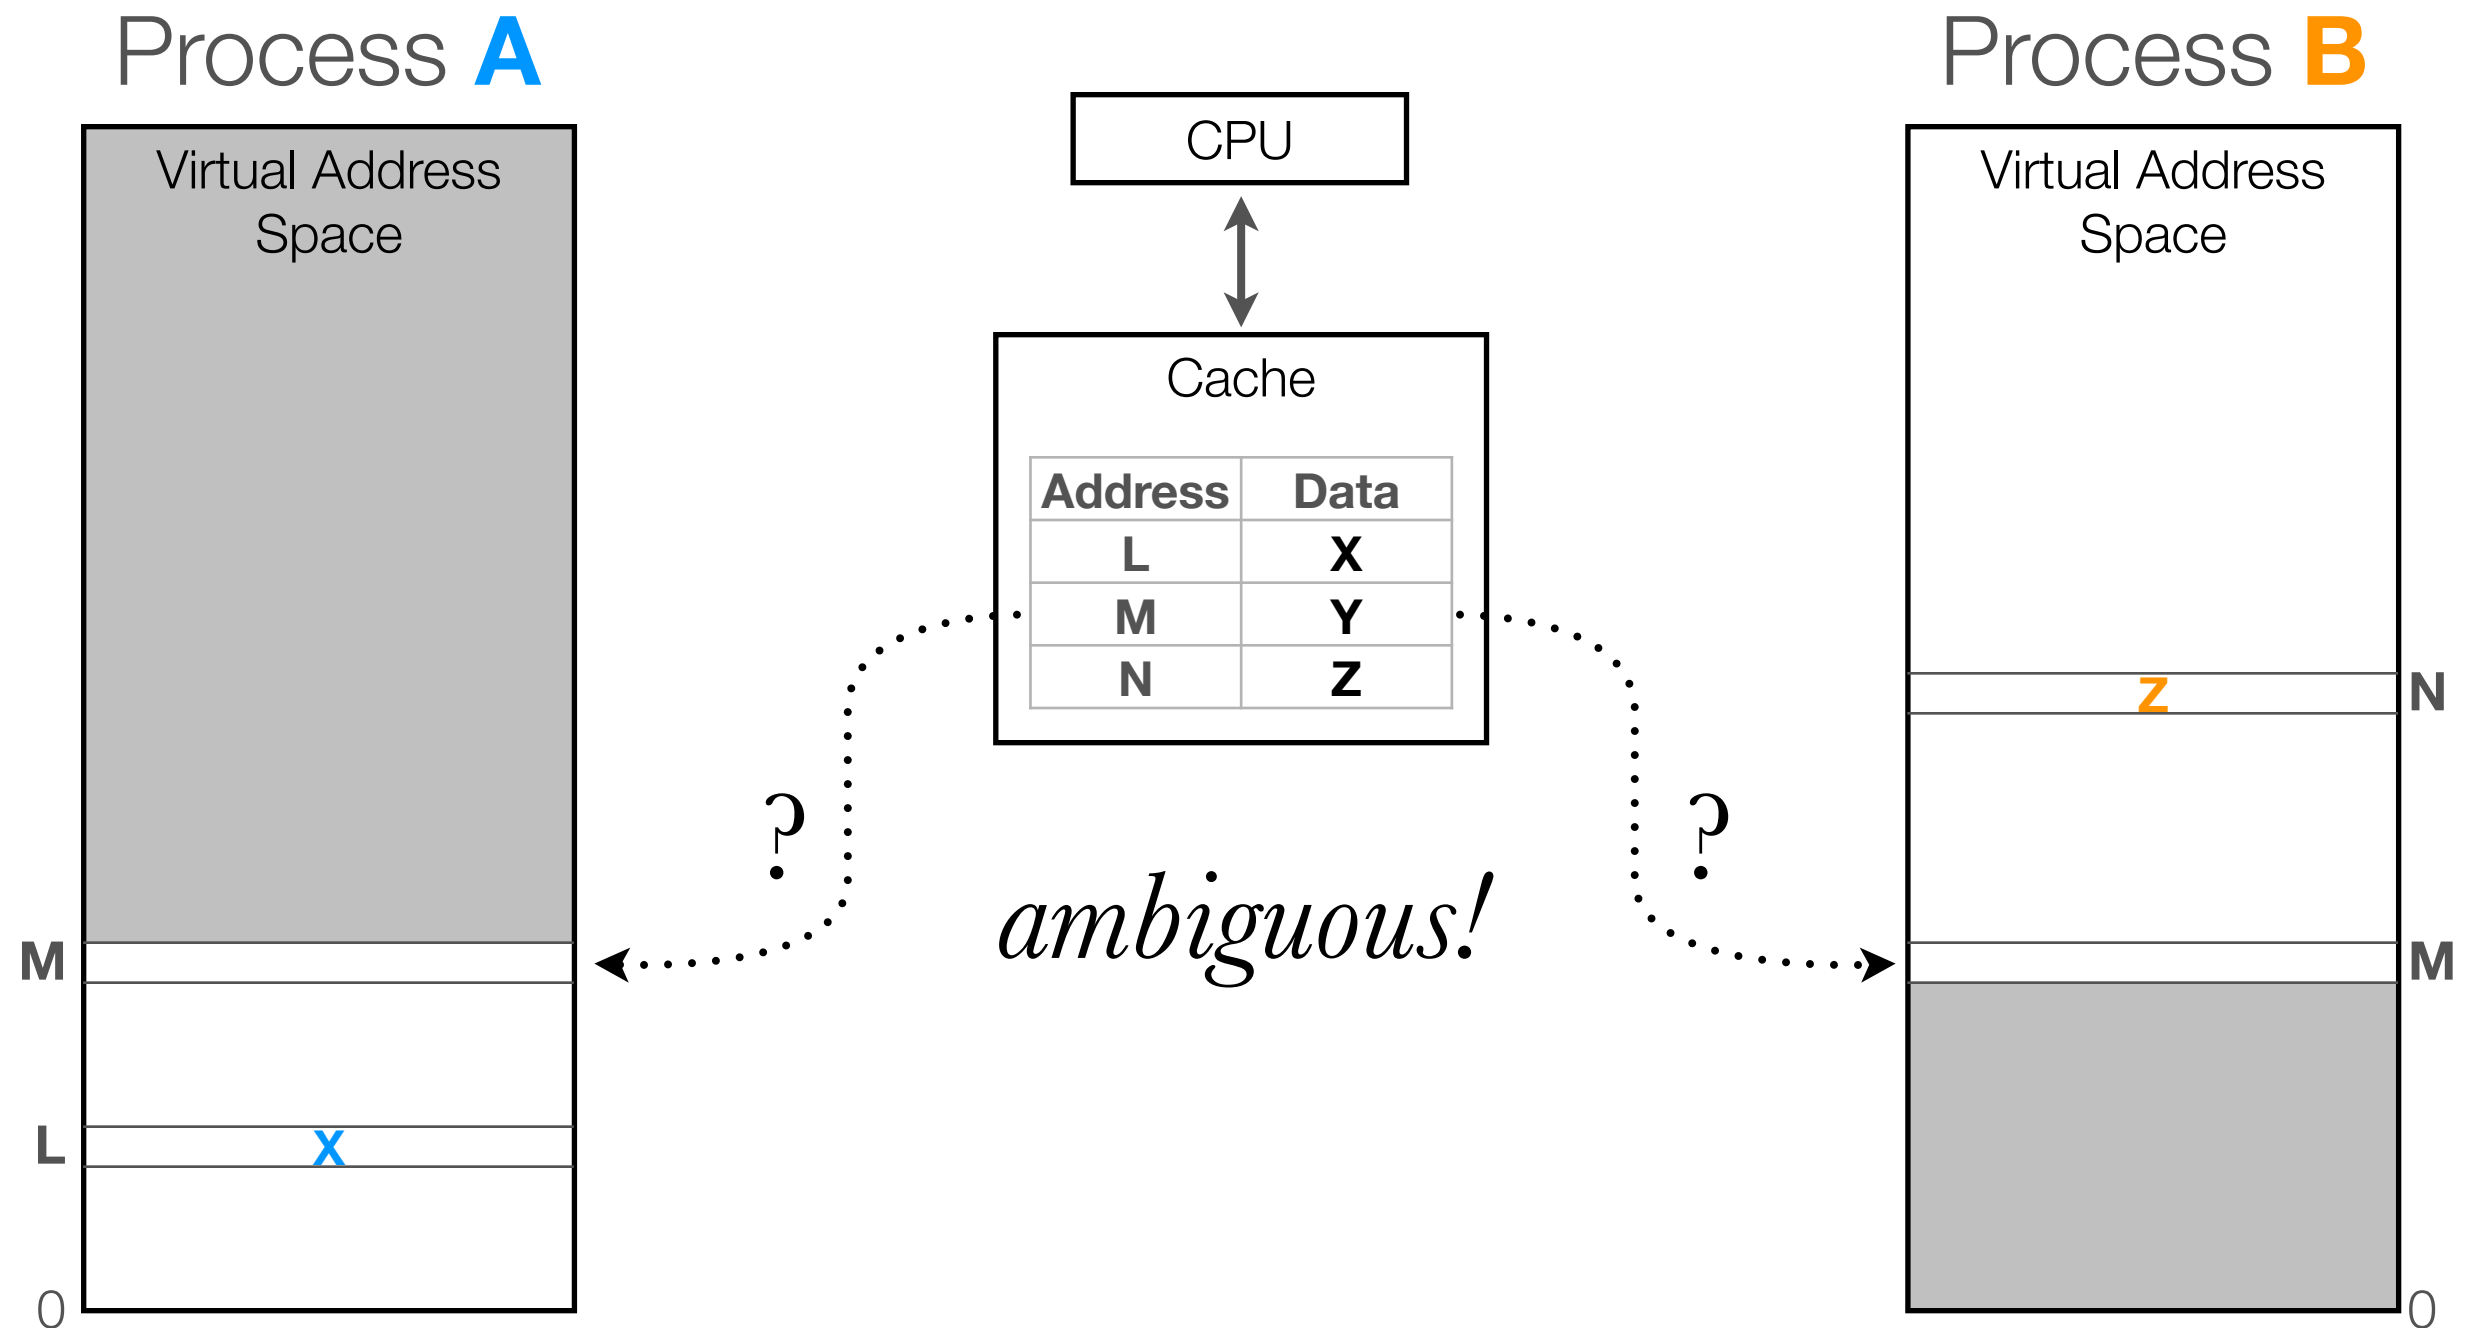
integration with caches?



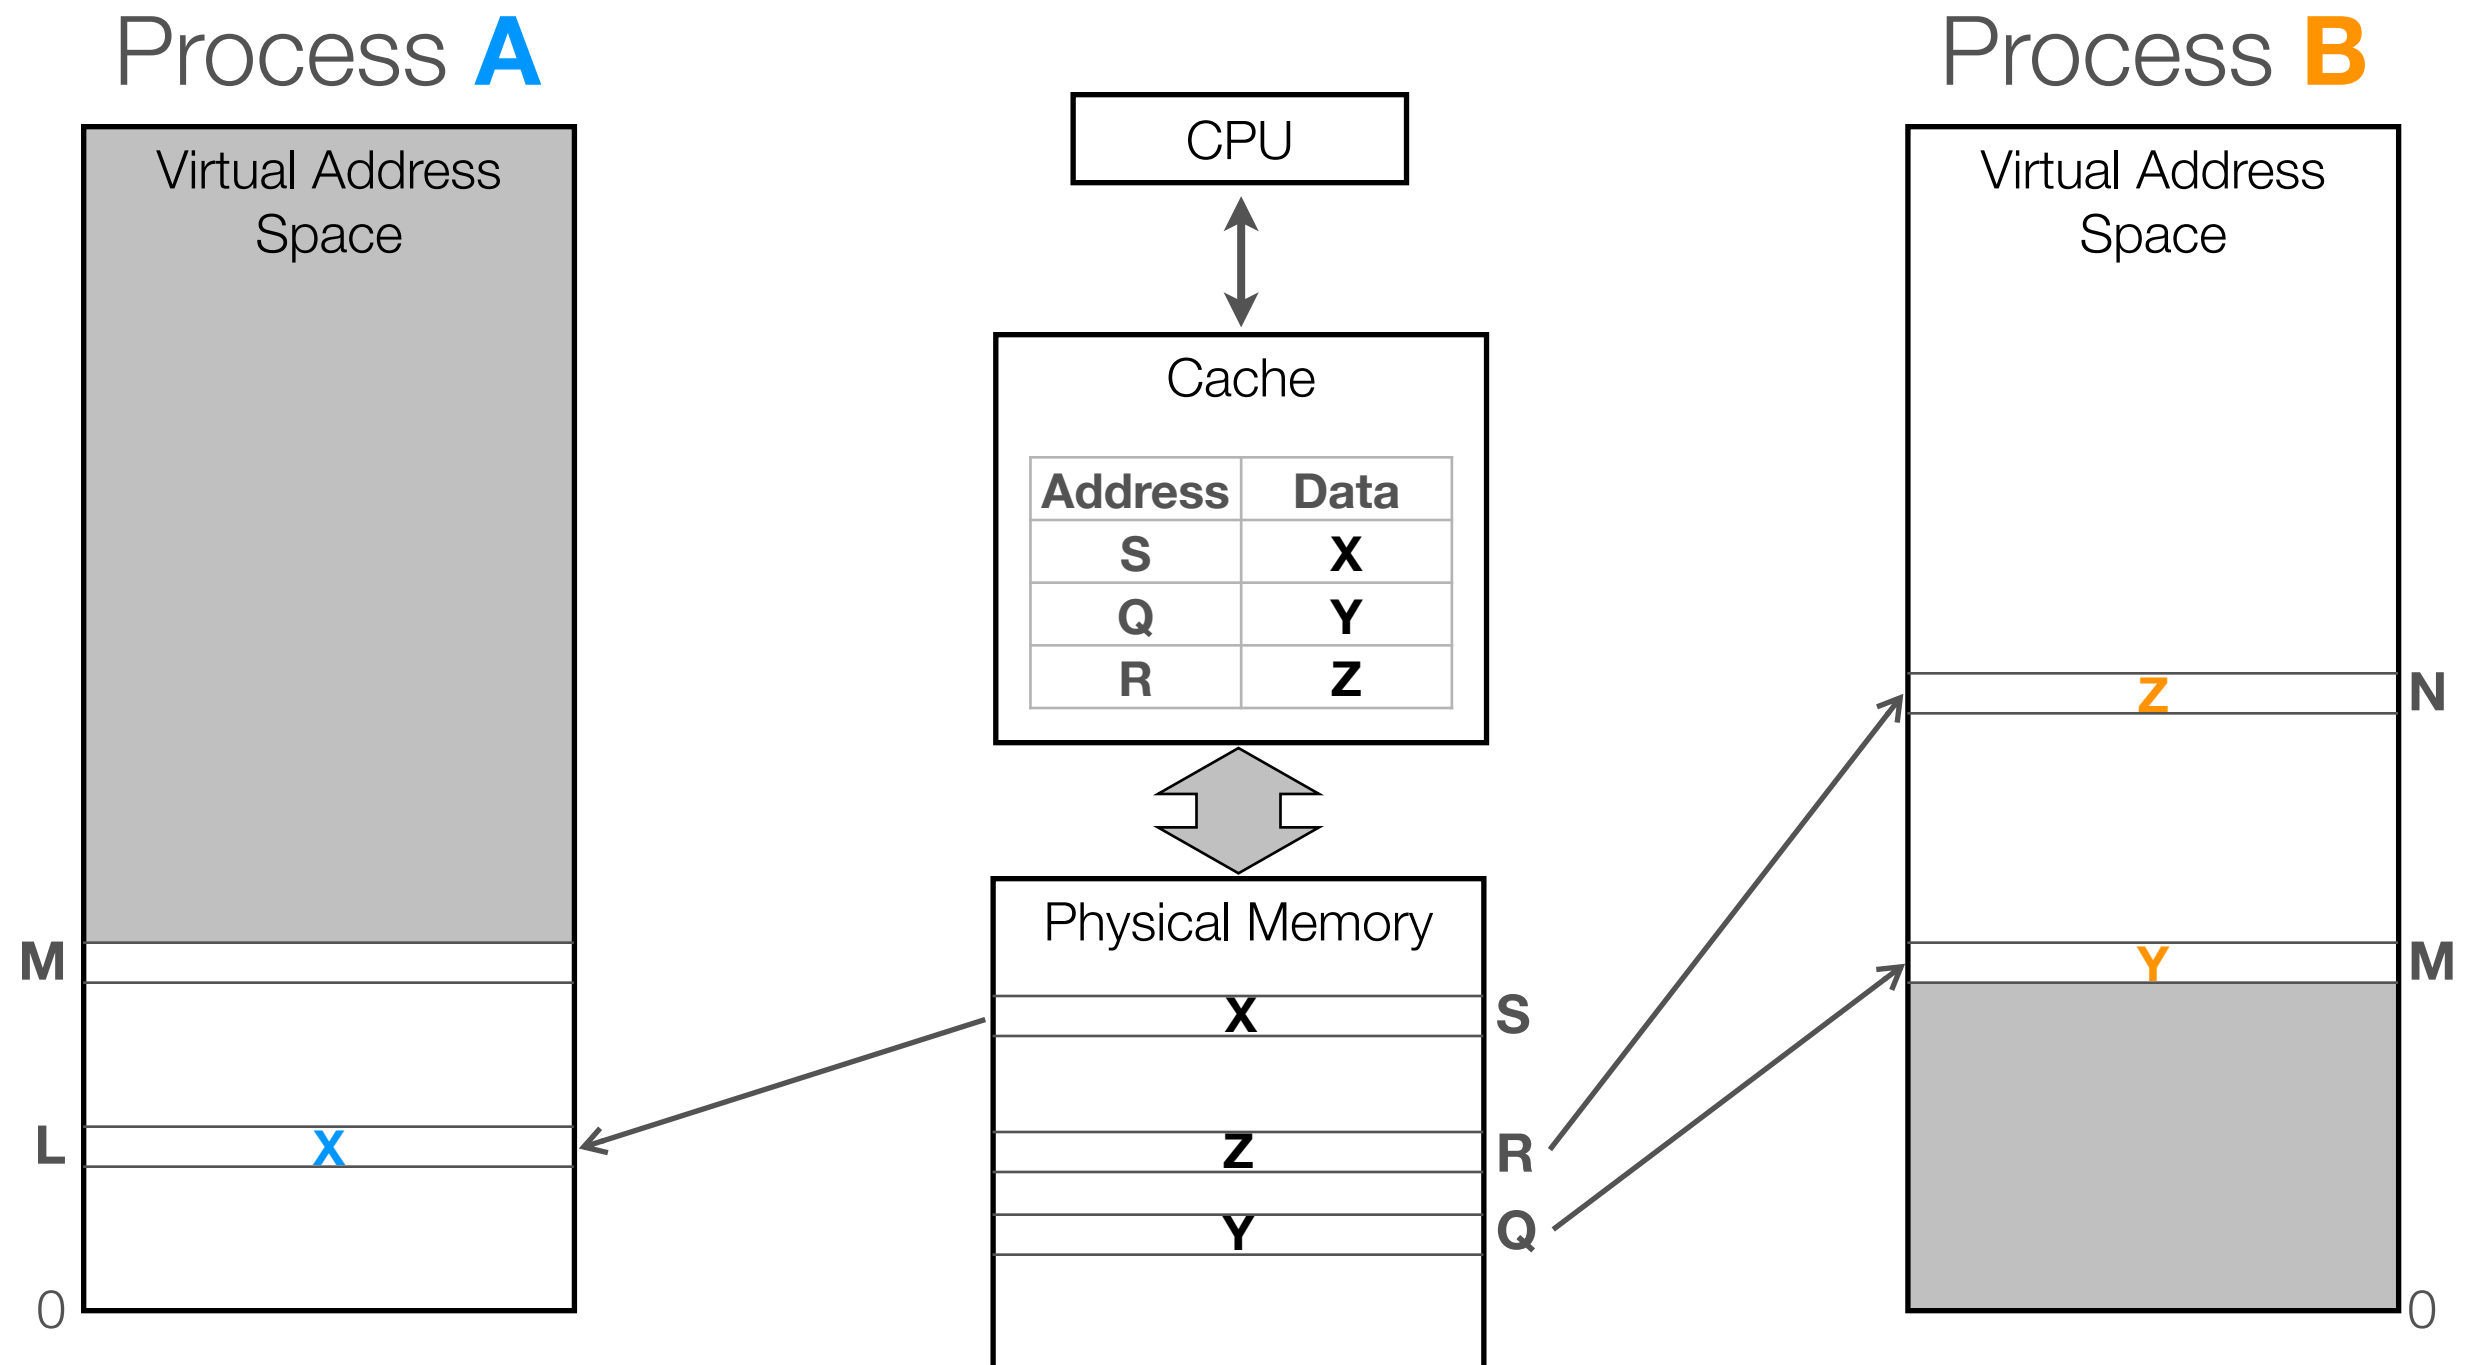
Q: do caches use physical or virtual addresses for lookups?



# Virtual address based Cache



# Physical address based Cache

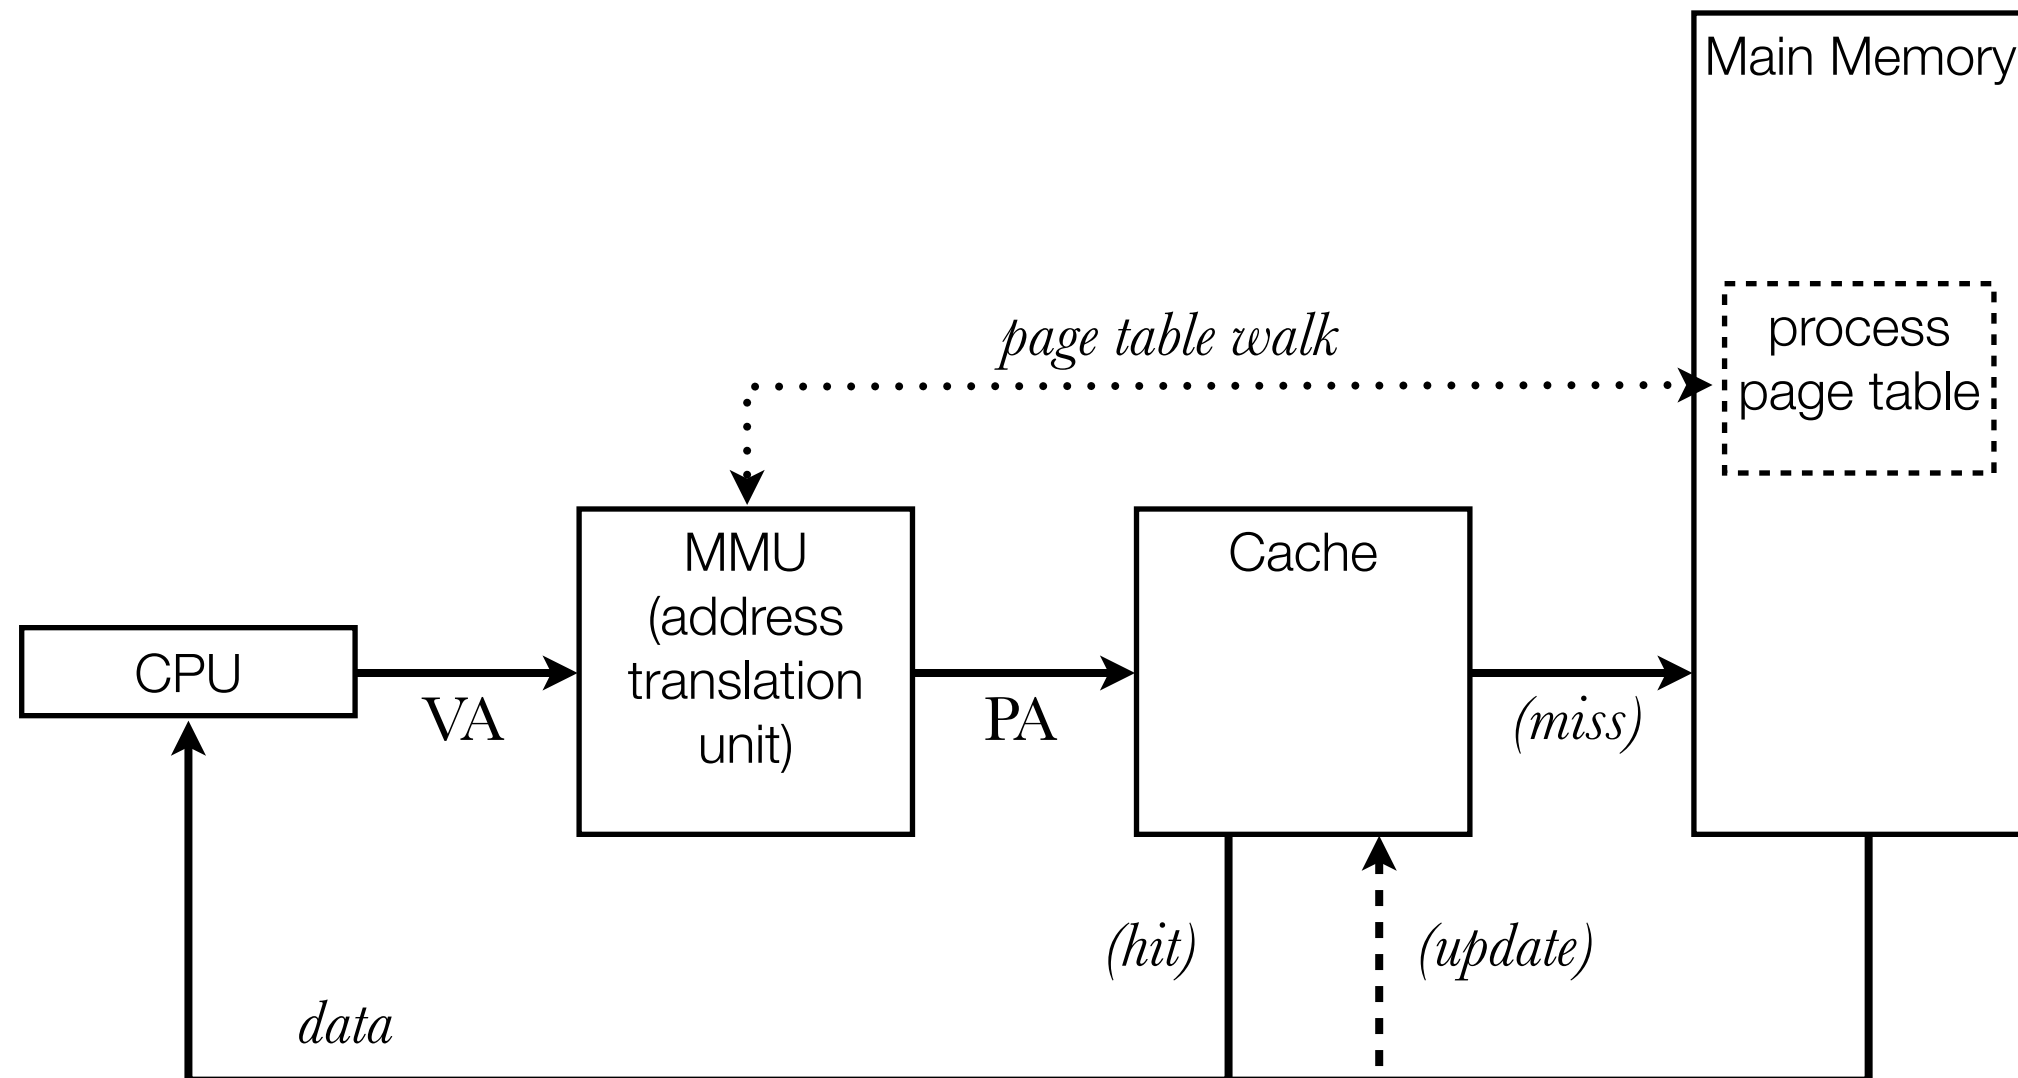




Q: do caches use physical or virtual addresses for lookups?

A: caches typically use *physical* addresses





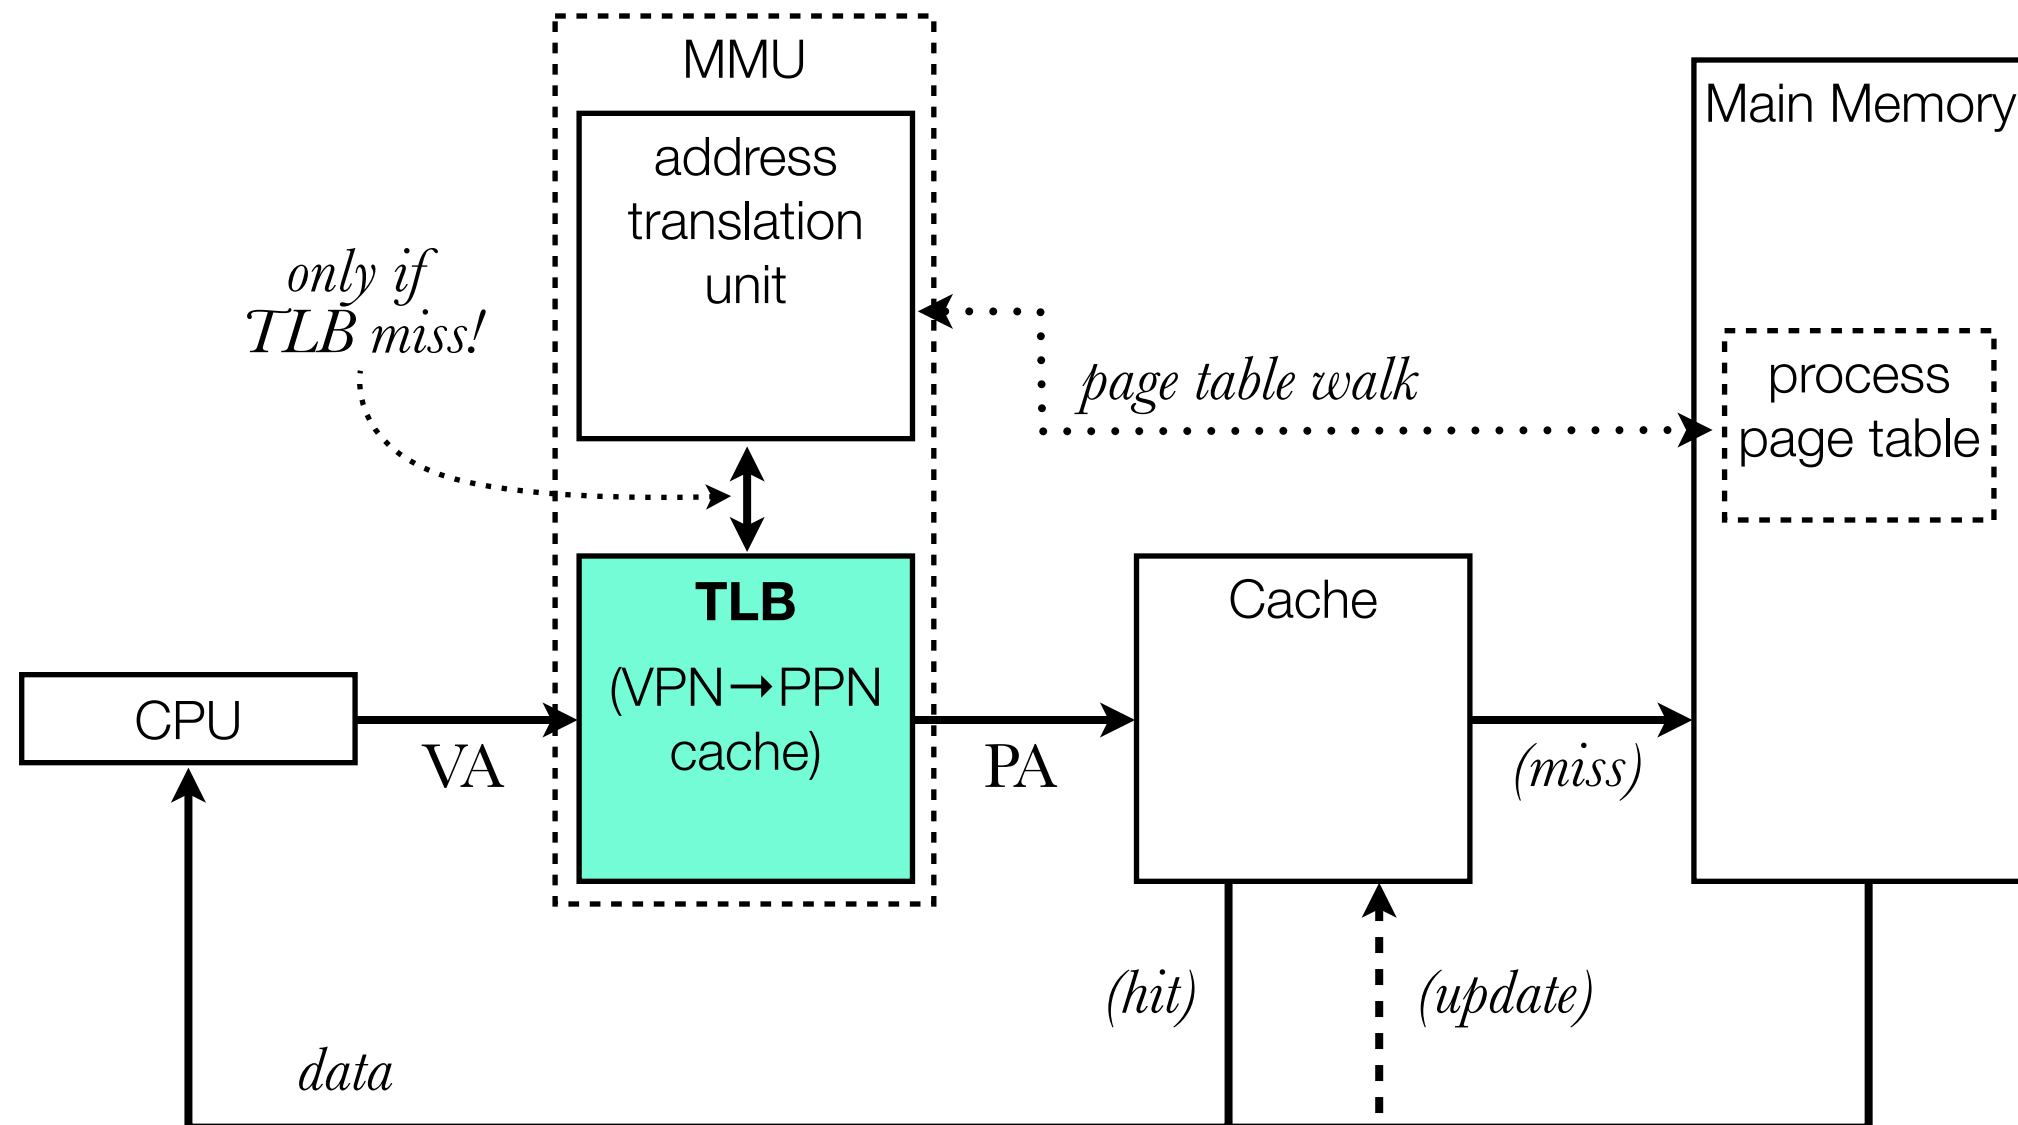
**%\*@&#!!!**



saved by hardware:

the *Translation Lookaside Buffer* (TLB) — a cache used solely for  $\text{VPN} \rightarrow \text{PPN}$  lookups

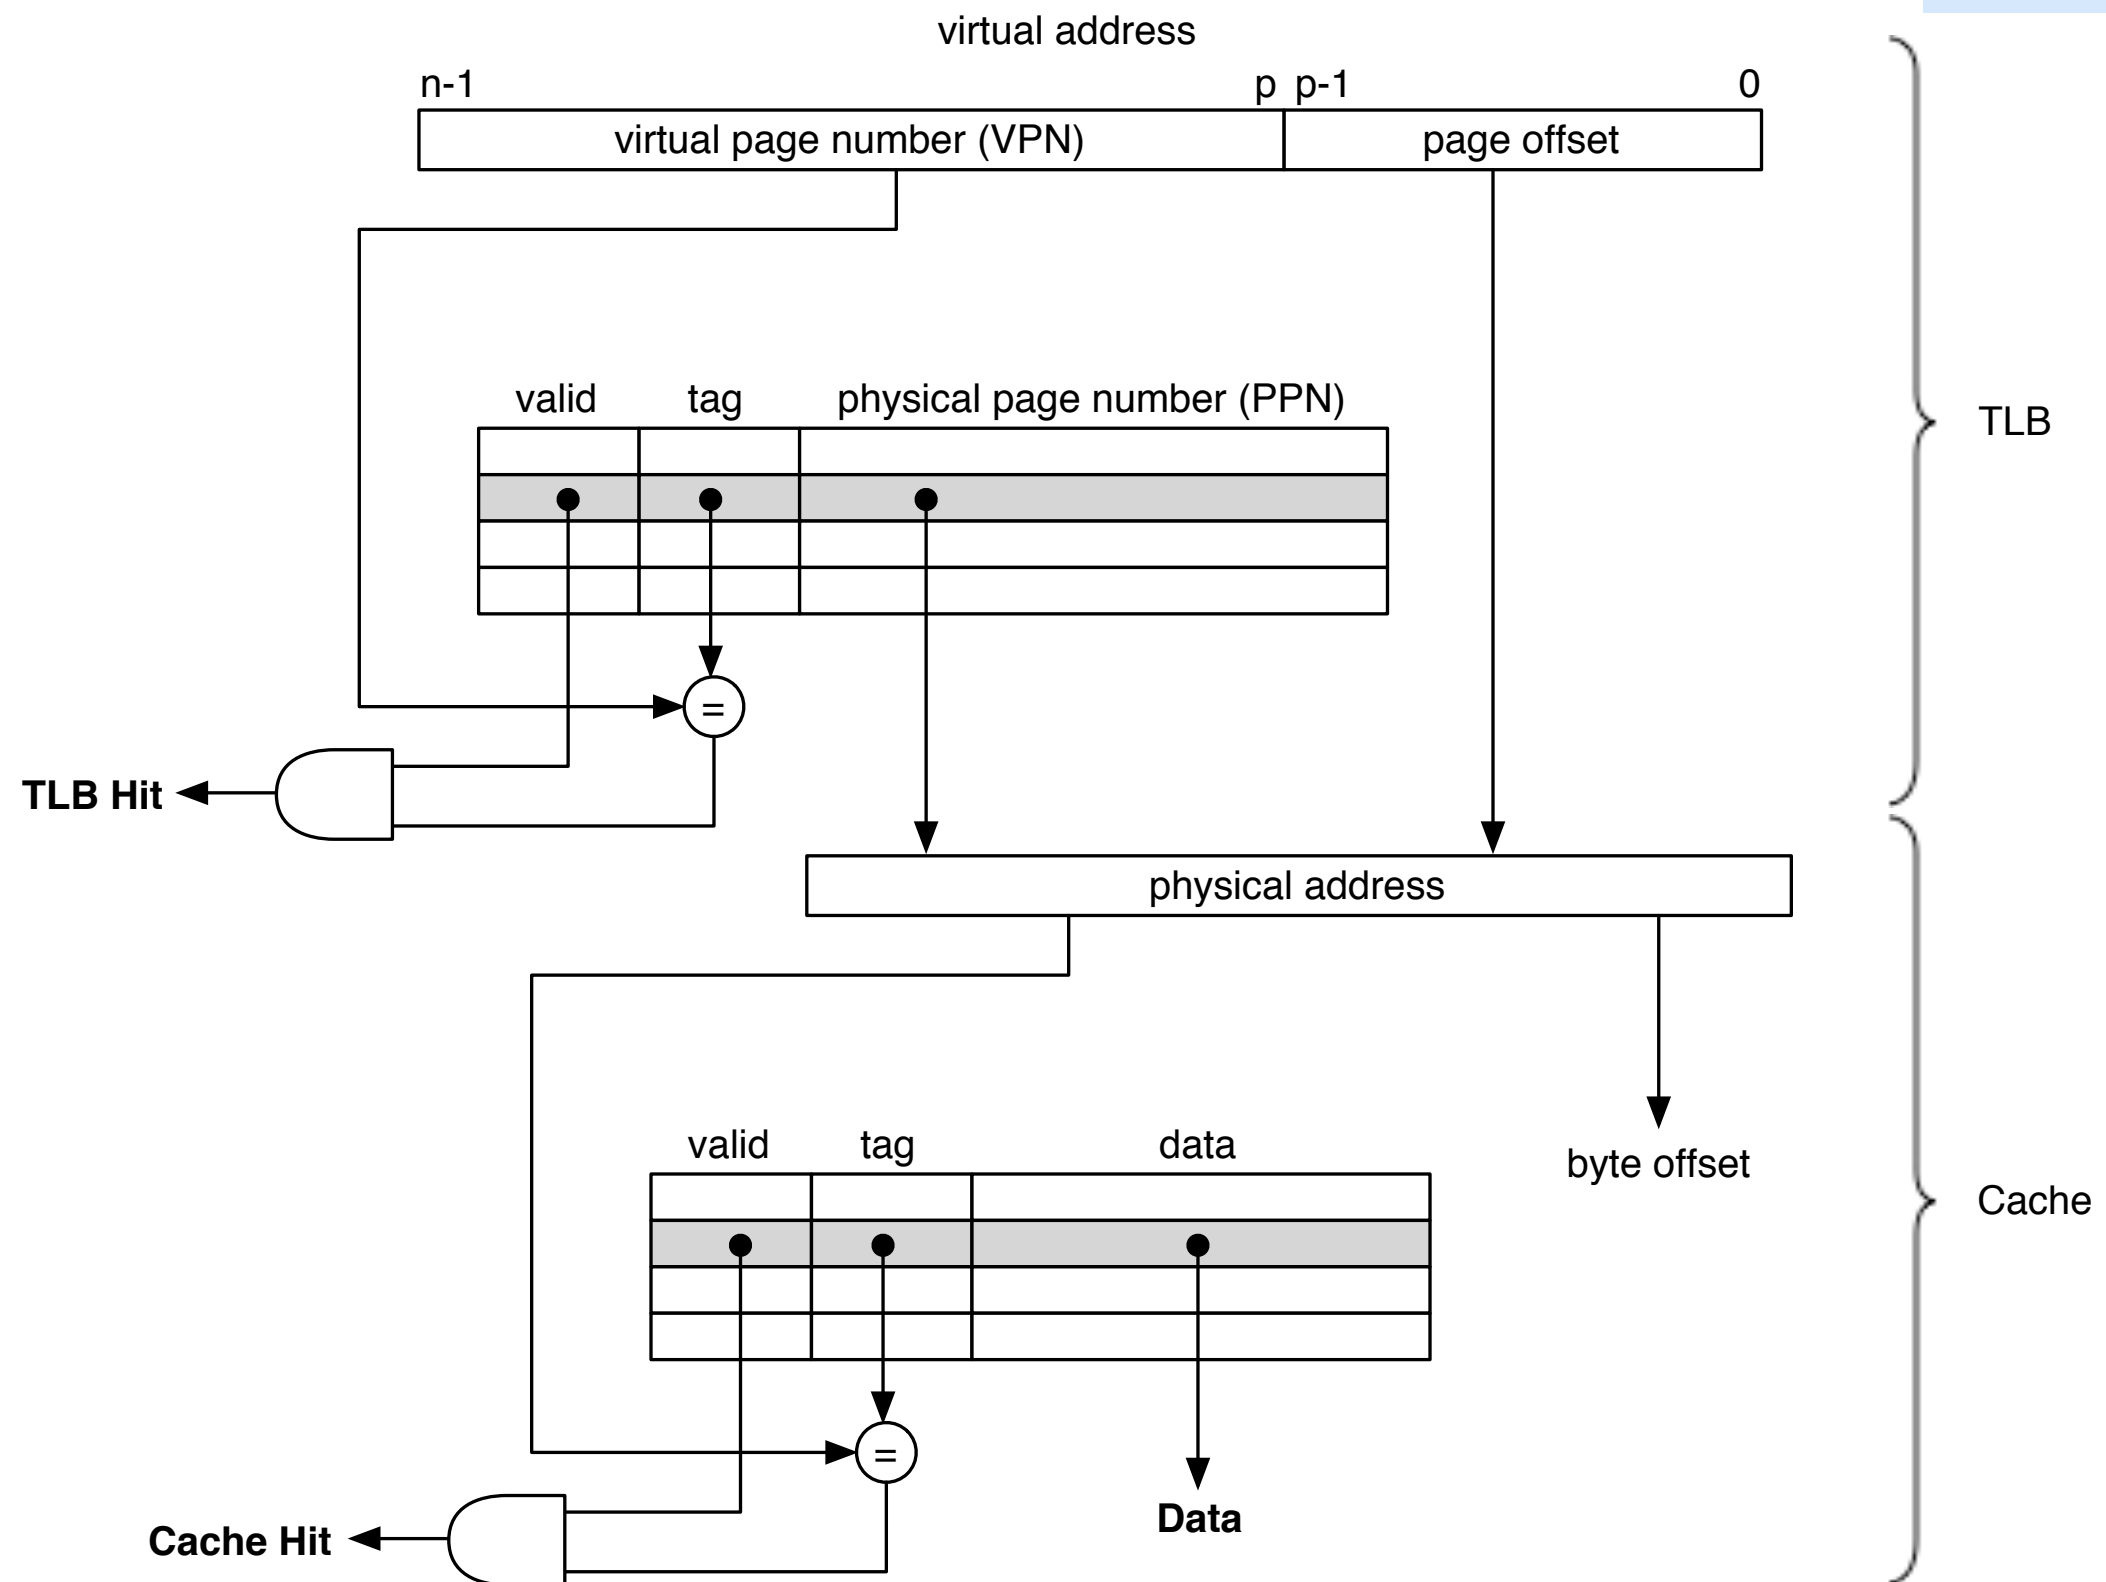




# TLB + Page table

(exercise for reader: revise earlier translation diagrams!)





TLB mappings are *process specific* —  
requires flush & reload on context switch

- some architectures store PID (aka  
“virtual space” ID) in TLB



Familiar caching problem:

- TLB caches a few thousand mappings
- vs. *millions* of virtual pages per process!



we can improve TLB hit rate by reducing  
the number of pages ...

by increasing the size of each page





compute # pages for 32-bit memory for:


- 1KB, 512KB, 4MB pages

- $2^{32} \div 2^{10} = 2^{22} = 4\text{M pages}$


- $2^{32} \div 2^{19} = 2^{13} = 8\text{K pages}$

- $2^{32} \div 2^{22} = 2^{10} = 1\text{K pages}$   
(not bad!)





Process **A**

## Physical Memory





lots of wasted space!

Process **B**

Process **A**

Physical Memory

Process **B**

increasing page size results in increased  
*internal fragmentation* and *lower utilization*



i.e., TLB effectiveness needs to be  
balanced against memory utilization



so what about 64-bit systems?

$2^{64} = 16$  Exabyte address space

$\approx 4$  billion x 4GB



most modern implementations support a  
max of  $2^{48}$  (256TB) addressable space



page table size (assuming 4K page size)?

- # pages  $= 2^{48} \div 2^{12} = 2^{36}$
- PTE size = 8 bytes (64 bits)
- PT size  $= 2^{36} \times 8 = 2^{39}$  bytes  
**= 512GB**





**512GB**

(just for the virtual memory *mapping* structure)

(and we need *one per process*)



(these things aren't going to fit in memory)





instead, use *multi-level* page tables:

- split an address translation into two (or more) separate table lookups
- unused parts of the table don't need to be in memory!



# “toy” memory system

- 8 bit addresses
- 32-byte pages



## Page Table

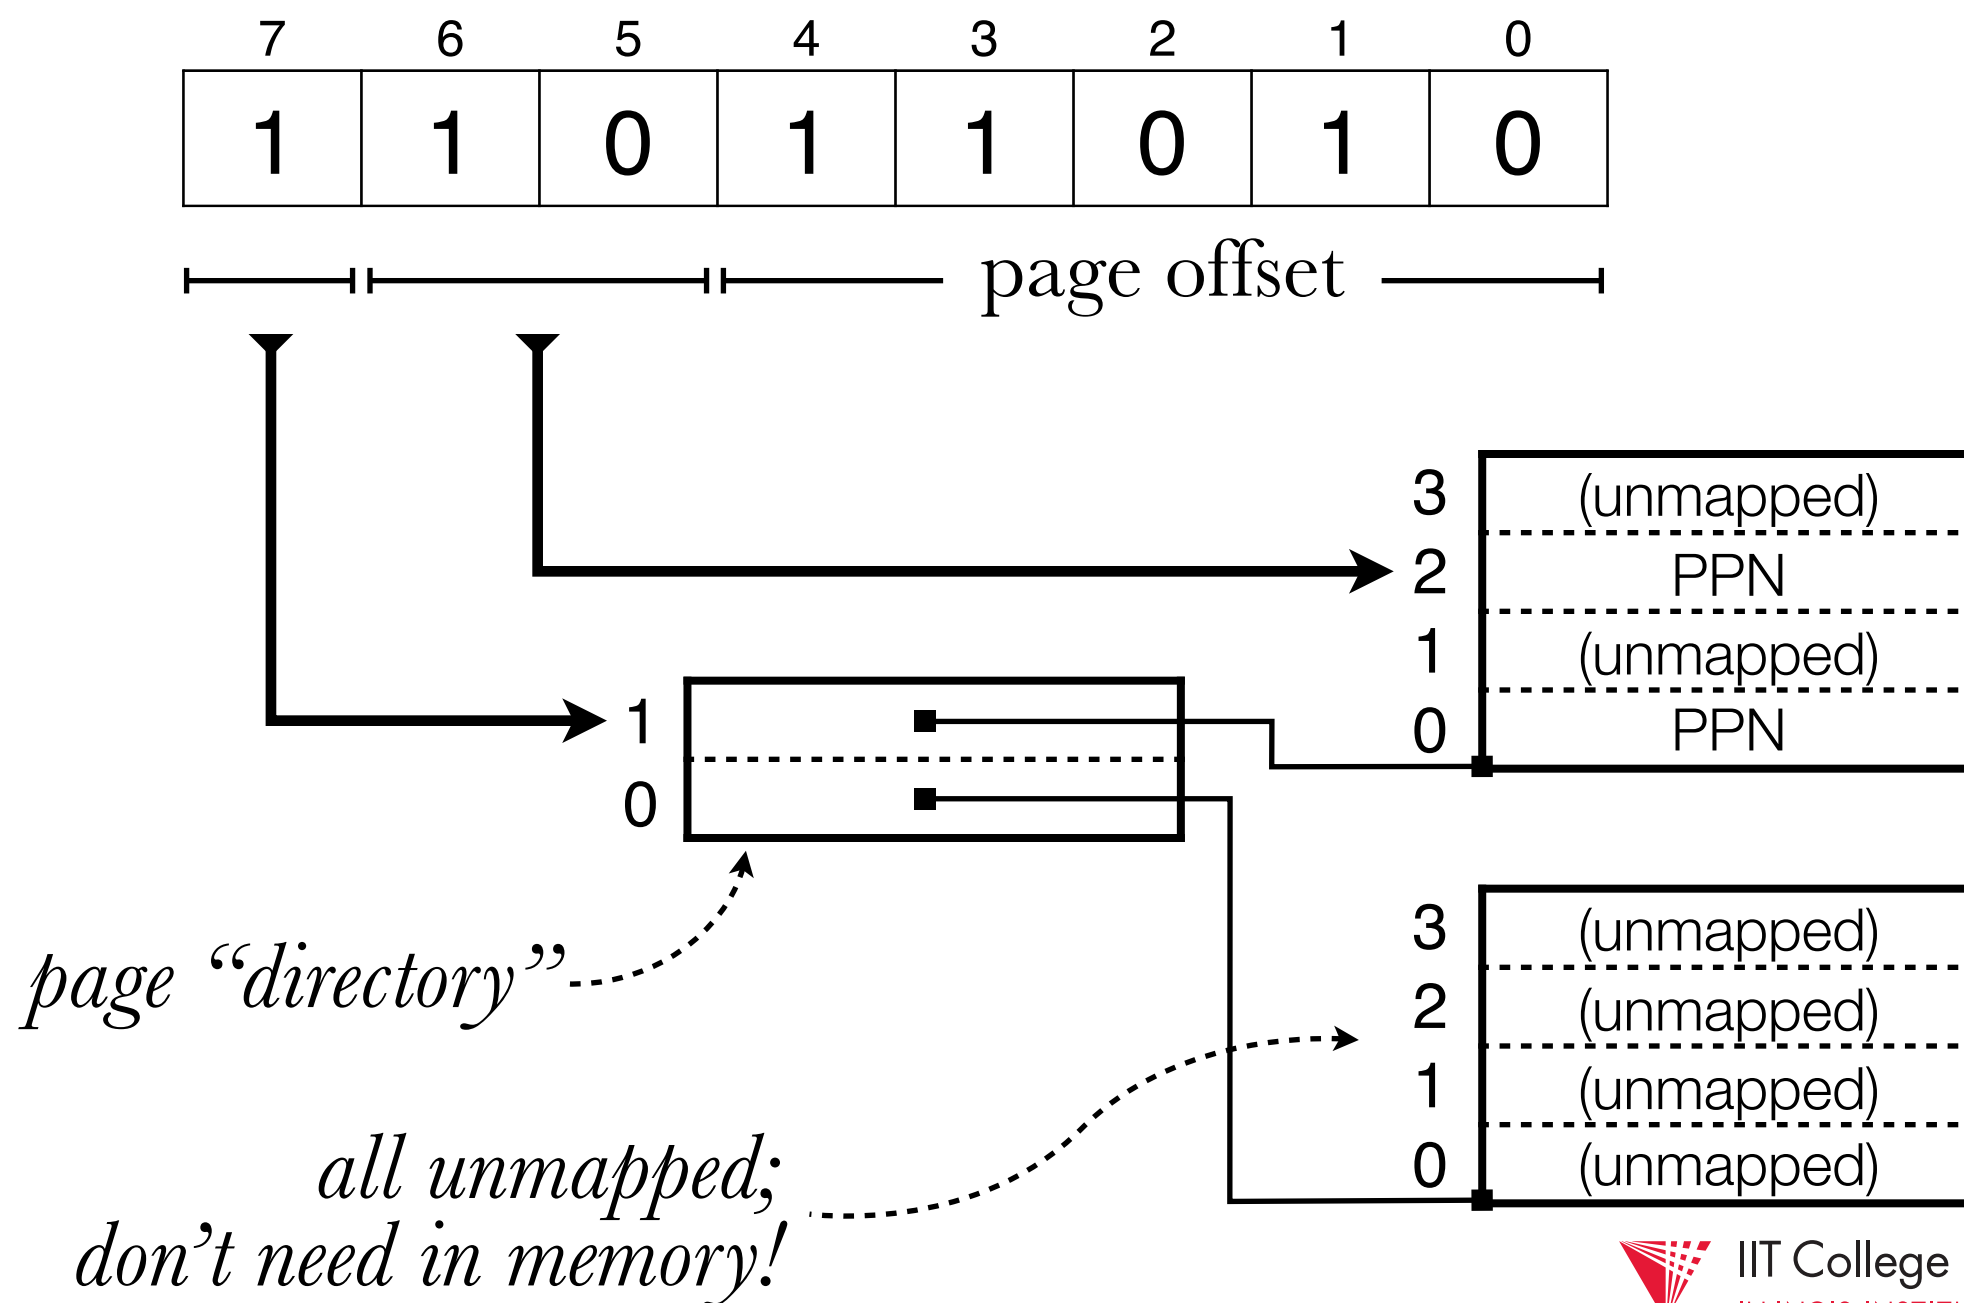
7	(unmapped)
6	PPN
5	(unmapped)
4	PPN
3	(unmapped)
2	(unmapped)
1	(unmapped)
0	(unmapped)

*all 8 PTEs  
must be in  
memory at  
all times*




# “toy” memory system

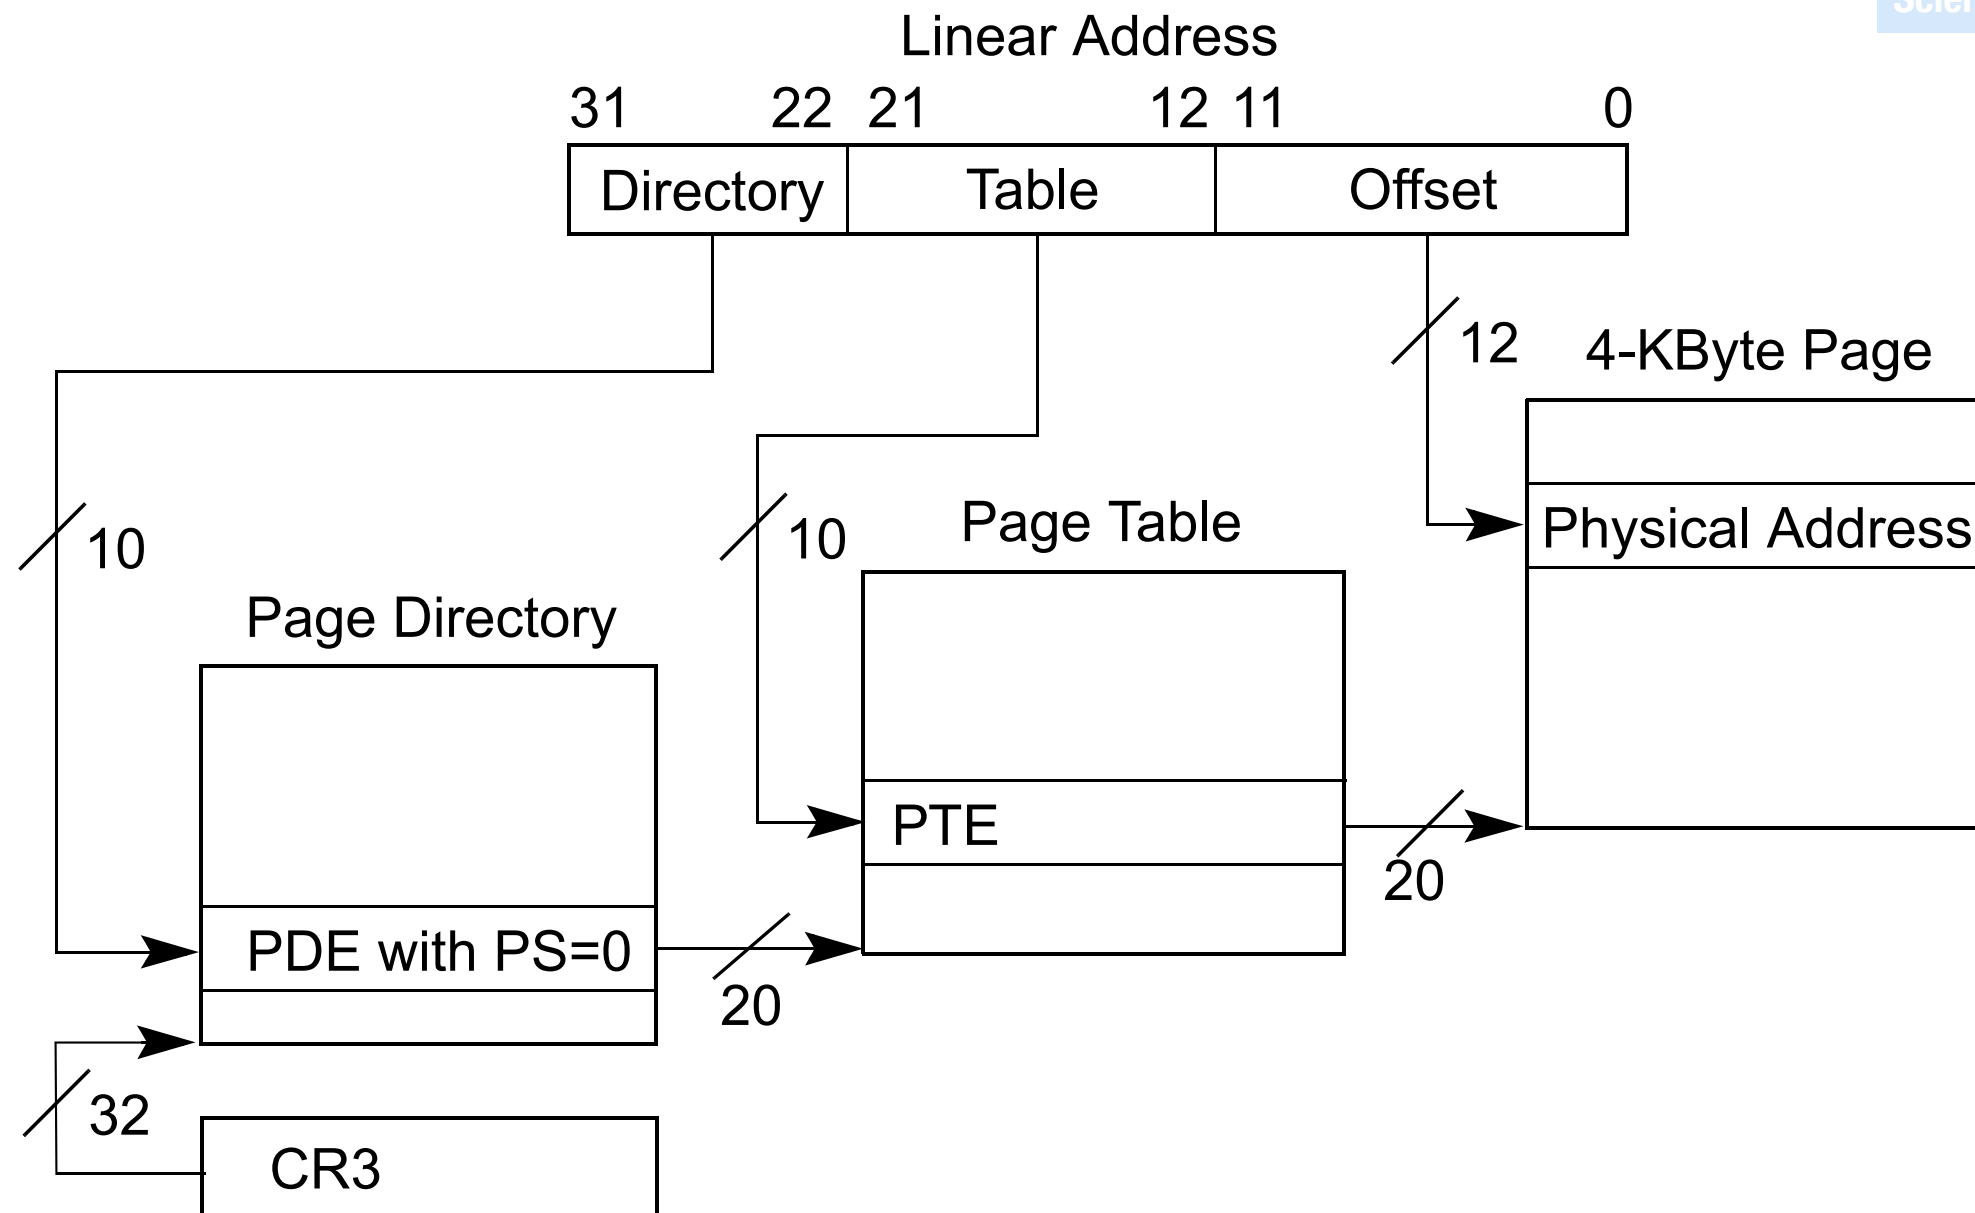
- 8 bit addresses
- 32-byte pages



# “toy” memory system

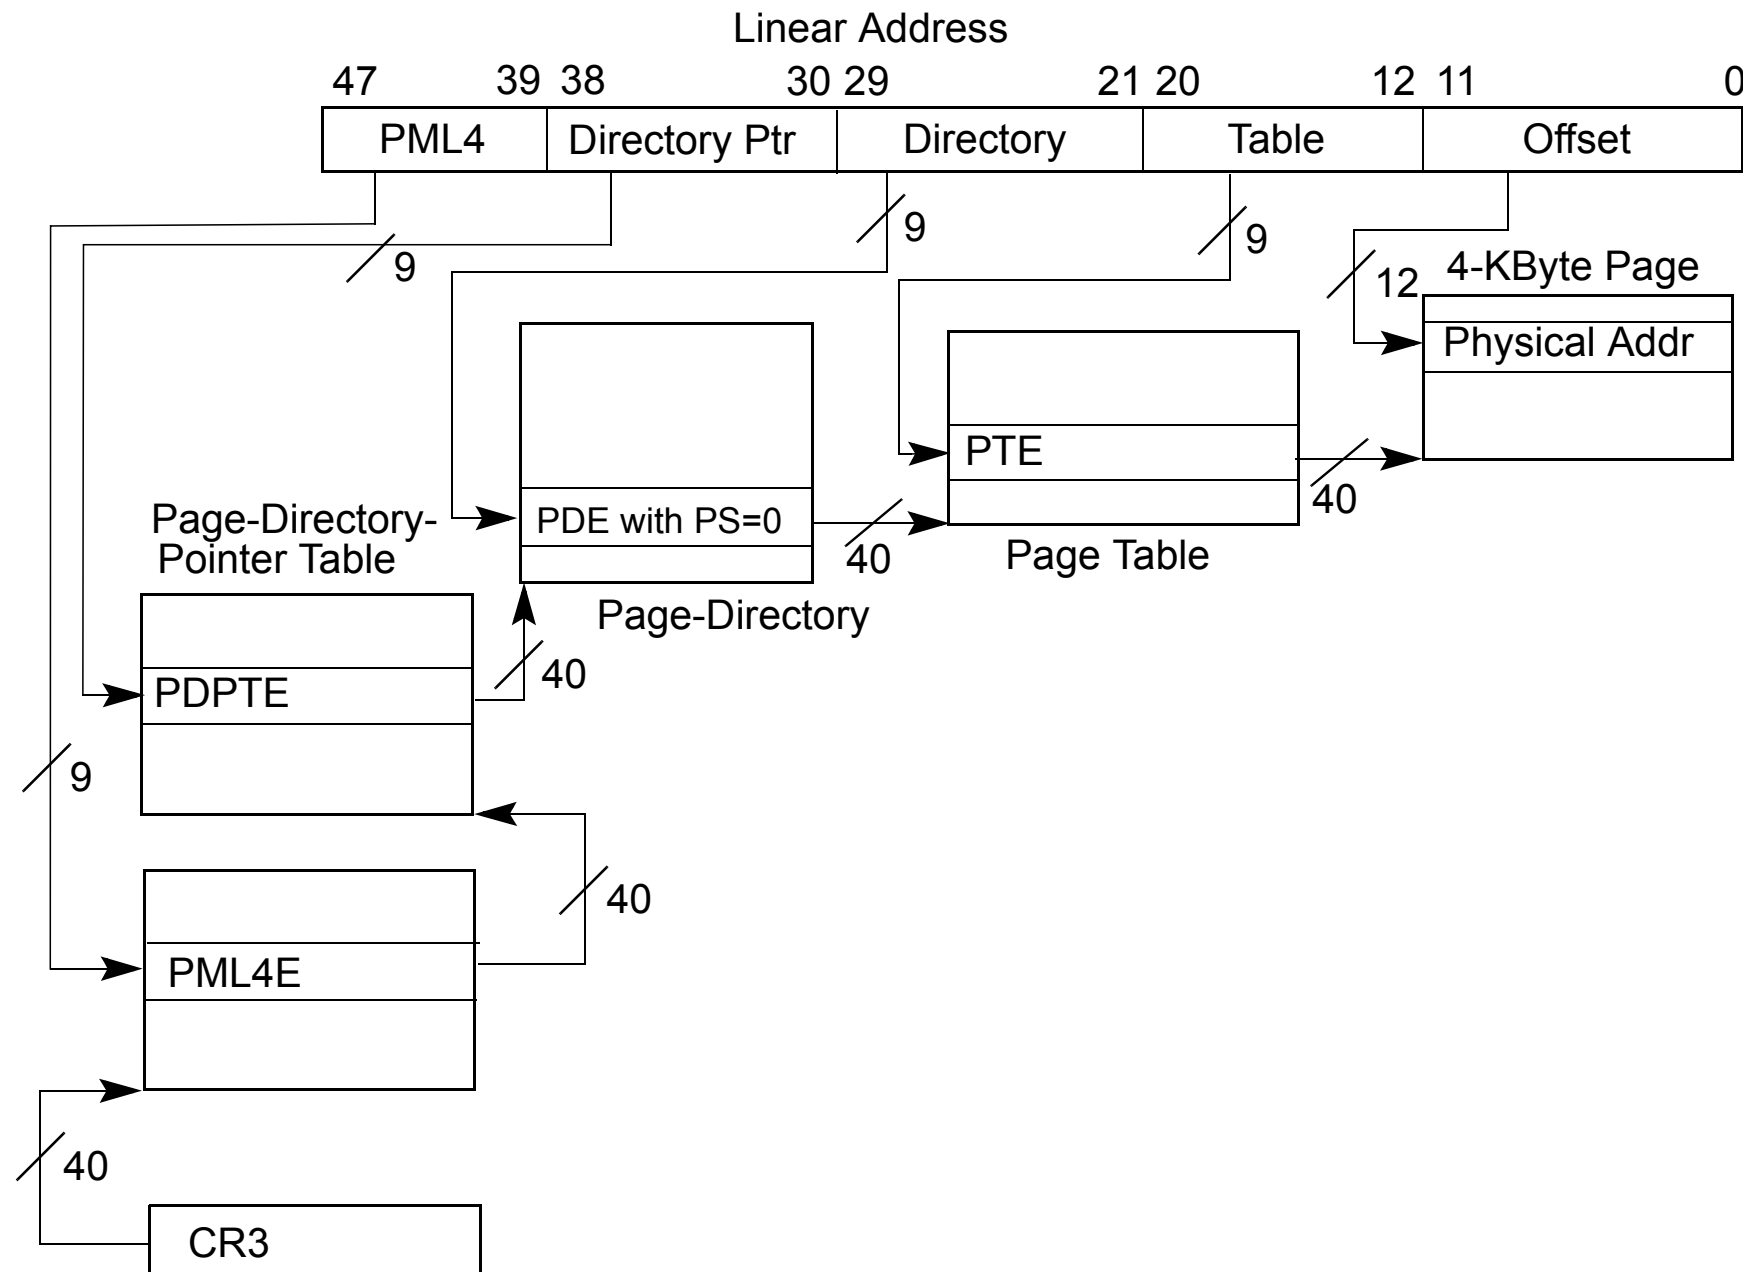
- 8 bit addresses
- 32-byte pages





# IA-32 paging (4KB pages)





# x86-64 paging (4KB pages)

